

# DYNAMO: Dynamic Skill-Tool Evolution for Vision-Language Agents

Anonymous ACL submission

## Abstract

Improving vision-language models (VLMs) on visual reasoning typically requires retraining or hand-designed prompts and tools. We present DYNAMO, a training-free framework that adapts a frozen VLM without any weight updates. On a small labeled training subset, the agent inspects its own correct and incorrect attempts and evolves two complementary capabilities: reusable reasoning skills for cognitive bottlenecks, and executable visual tools for perceptual ones. Each generated tool is paired with a skill that specifies when to invoke it, and both capability types accumulate in a persistent library. Across four visual reasoning benchmarks and five VLM backbones, DYNAMO improves direct inference on all 20 model–benchmark settings (avg. +5.6 acc). When the tool set is given in advance, the framework learns when to call each tool, and per-step tool choice improves on every tested backbone. Against task-specific RL (VTool-RL, DeepEyes), DYNAMO closes 65–99% of the RL gap at a fraction of the compute, and combines additively with RL when available.

## 1 Introduction

Vision-language models (VLMs) have improved rapidly, yet adapting a VLM to a new visual-reasoning task family still typically requires manually curated SFT data or a hand-designed RL pipeline (Li et al., 2025). We ask whether the agent can build its own task-specific capability set instead, by inspecting its own behaviour on small training subsets and without weight updates. Figure 1 contrasts this capability-evolution route with the typical per-task SFT/RL adaptation pipeline.

We instantiate this idea in DYNAMO, a training-free framework that evolves two complementary capability types from a small labeled training subset and dynamically equips them at inference. **Skills** are structured Markdown SOPs for *cognitive* bottlenecks, where the visual evidence is available but



Figure 1: **Hand-crafted SFT/RL vs. DYNAMO.** Top: each new task family requires hand-curated SFT data and a custom RL pipeline. Bottom: DYNAMO instead evolves a task-family-specific skill and tool set from a small labeled training subset, with the VLM frozen.

the agent’s reasoning procedure is weak. **Tools** are short Python programs for *perceptual* bottlenecks, where the agent needs a transformed view of the input before the evidence becomes legible: a crop or zoom into a region of interest, a contrast adjustment, a chart re-rendering with enlarged axis labels, the extraction of a specific color or chart layer, or a saliency-guided sub-region extraction. On each iteration over a sampled sub-training set, DYNAMO diagnoses both the correct and the incorrect attempts using their images and reasoning traces, proposes multiple candidate skill+tool combinations, and promotes the highest-validation-accuracy candidate into a persistent library. A mastery phase further learns when each tool is safe to

invoke, so that promoted capabilities are deployed selectively rather than indiscriminately.

We evaluate three questions. **(I) Does DYNAMO work across diverse benchmarks and backbones?** On ChartQA (Masry et al., 2022), MathVista (Lu et al., 2024), HRBench4K (Wang et al., 2025), and V\* (Wu and Xie, 2024) across five VLM backbones, skill-tool co-evolution improves direct inference on all 20 model–benchmark settings, averaging +5.6 accuracy points. **(II) Does the framework extend to a tool-mastery setting with a curated tool set?** On GTA (Wang et al., 2024b), learning when to call each pre-provided tool improves step-level tool selection, argument prediction, and instruction-following accuracy across all tested backbones. **(III) Can DYNAMO match task-specific RL across different task families?** We compare DYNAMO to a representative RL method on two task families: structured visual QA (VTool-R1 (Wu et al., 2025a)), and high-resolution perception (DeepEyes (Zheng et al., 2025)). DYNAMO matches or approaches RL accuracy at a fraction of the compute (65–99% gap recovery on chart/table; 91.1 vs. 90.1 on V\*-7B perception), and combines additively with RL when both are available.

## Contributions.

- (1) A **problem formulation** that recasts per-task VLM adaptation as capability evolution: a frozen agent builds its own skill and tool library from a small labeled training subset.
- (2) A **multi-candidate evolution loop** that diagnoses correct and incorrect attempts, proposes candidate skill+tool combinations, promotes the best by validation accuracy, learns tools’ applicability, and grows the library iteratively.
- (3) **Empirical evidence** across four visual-reasoning benchmarks on five VLM backbones, the GTA tool-mastery benchmark, and two RL-comparison task families, with additive composition when RL is also available.

## 2 Related Work

**Self-improving agents.** A growing line of work lets LLM agents improve from experience without gradient updates: Reflexion (Shinn et al., 2023) keeps a verbal reflection buffer that resets between episodes; Voyager (Wang et al., 2024a) and ExpeL (Zhao et al., 2024) accumulate persistent skill libraries in game and tool-use domains; AutoManual (Zhao et al., 2024) and EvolveR (Wu

et al., 2025b) refine instruction manuals or reasoning strategies via self-play; Trace2Skill (Ni et al., 2026) distills trajectory-local lessons into transferable agent skills; and a broader line uses self-distillation to drive self-evolution (Zhang et al., 2026; Sun et al., 2025; Xu et al., 2025).

**Tool creation for language models.** A complementary line generates new tools on demand: CREATOR (Qian et al., 2023) prompts an LLM to write Python helpers for problems it cannot solve directly; LATM (Cai et al., 2024) and ToolLLM (Qin et al., 2024) scale this to large API collections. The generated tools are typically text-domain utilities, such as arithmetic helpers, unit converters, and API wrappers, evaluated on math word problems and similar text-symbolic benchmarks.

**Visual reasoning with tools.** A growing line equips VLMs with visual processing tools. VTool-R1 (Wu et al., 2025a), DeepEyes (Zheng et al., 2025; Hong et al., 2025), PixelReasoner (Su et al., 2025), and V-Thinker (Qiao et al., 2025) rely on SFT and/or RL with curated trajectories to teach VLMs to invoke a fixed tool inventory. ZoomEye (Shen et al., 2025) and ReFocus (Fu et al., 2025) are training-free but assume a fixed zoom/search or editing interface; Lever LM (Yang et al., 2024) configures in-context example sequences to leverage VLMs without re-training. A parallel program-synthesis line composes hand-curated tool libraries of detectors, OCR, and arithmetic modules (VisProg (Gupta and Kembhavi, 2023), ViperGPT (Surís et al., 2023), Chameleon (Lu et al., 2023)); MMR-Bench (Ma et al., 2026) evaluates such routing across diverse backbones and tools, and EcoAlign (Cheng et al., 2025) frames the cost of adapting VLMs.

## 3 Method

### 3.1 Problem Formulation

Let  $\mathcal{D}_{\text{train}} = \{(x_i, y_i, \mathbf{v}_i)\}_{i=1}^k$  denote a small training subset of  $k$  cases, where  $x_i$  is the question,  $y_i$  the ground-truth answer, and  $\mathbf{v}_i$  the associated visual input (one or more images). Let  $\pi_\theta$  be a frozen VLM backbone. We define a capability set  $\mathcal{C} = (\mathcal{S}, \mathcal{T})$ , where  $\mathcal{S}$  is a library of *skills* (structured reasoning SOPs) and  $\mathcal{T}$  is a library of *tools* (executable Python programs). The agent decomposes solving into a retrieval step and a reasoning




Figure 2: **DYNAMO evolution loop.** On a sampled sub-training set, DYNAMO diagnoses both correct and incorrect attempts, explores candidate skill+tool capabilities, and promotes the best by validation into a persistent library.

step:

$$f_{\mathcal{C}}(x, \mathbf{v}; \pi_{\theta}) = \pi_{\theta}(\cdot | x, \mathbf{v}, \text{Retrieve}(x, \mathbf{v}; \mathcal{C})), \quad (1)$$

where `Retrieve` returns the subset of skills and tools relevant to the current input. Given a held-out validation set  $\mathcal{D}_{\text{val}}$ , our goal is to learn  $\mathcal{C}$  that maximises agent accuracy without any weight update:

$$\mathcal{C}^* = \arg \max_{\mathcal{C}} \text{Acc}(f_{\mathcal{C}}; \mathcal{D}_{\text{val}}), \quad (2)$$

where the maximisation is over  $\mathcal{C}$  only and  $\pi_{\theta}$  remains frozen ( $\nabla_{\theta} = 0$ ).

### 3.2 Evolution Loop

Each iteration has three phases (Figure 2).

**Diagnose.** The agent (Eq. 1) attempts each case in  $\mathcal{D}_{\text{train}}$  once with the current  $\mathcal{C}$ , producing one reasoning trace  $\tau_i$  per case. We then sample  $\mathcal{D}_{\text{sub}} \subseteq \mathcal{D}_{\text{train}}$  across questions (not stratified per question), including both correct and incorrect attempts when both are present. If all attempts are correct, the iteration is skipped (no bottleneck to fix); if all are incorrect, the ANALYZERDECIDER proceeds on failures alone. The ANALYZERDECIDER inspects  $\mathcal{D}_{\text{sub}}$ —the questions  $x_i$ , ground-truth answers  $y_i$ , original images  $\mathbf{v}_i$ , reasoning traces  $\tau_i$ , and any intermediate tool outputs—and emits a root-cause analysis grounded in visual evidence, an identified bottleneck type (*cognitive*

when the visual evidence is available but the reasoning procedure is weak, or *perceptual* when the agent needs a transformed visual input), and an action  $a \in \{\text{skill}, \text{tool}, \text{both}\}$ . Letting the ANALYZERDECIDER choose freely from this three-action set on each case is the default DYNAMO behaviour (denoted *Full* in Section 4.2); restricting  $a$  to  $\{\text{skill}\}$  or  $\{\text{tool}\}$  yields the *Skill Only* and *Tool Only* ablations. Visual input here is essential: deciding whether a crop is misaligned, whether labels are legible, or whether a processed image introduced artefacts requires inspecting the image, not just the reasoning trace.

**Explore.** Conditioned on the diagnosis, the GENERATOR proposes  $M$  candidate skill+tool combinations, each addressing the diagnosed bottleneck with a different strategy or implementation.

**Validate and promote.** Writing each candidate as  $c_m = (s_m, t_m)$ , we evaluate every  $c_m$  on  $\mathcal{D}_{\text{train}}$  and promote the highest-accuracy one:

$$c^* = \arg \max_{c \in \{c_1, \dots, c_M\}} \text{Acc}(f_{c \cup c}; \mathcal{D}_{\text{train}}),$$

$$\mathcal{C} \leftarrow \begin{cases} \mathcal{C} \cup \{c^*\} & \text{if } \text{Acc}(f_{c \cup c^*}) > \text{Acc}(f_c), \\ \mathcal{C} & \text{otherwise,} \end{cases} \quad (3)$$

where both accuracies are computed on  $\mathcal{D}_{\text{train}}$ . The guard ensures monotonic training-set accuracy: an iteration in which no candidate strictly improves over the current  $f_c$  leaves  $\mathcal{C}$  unchanged. Eq. 3 is

a training-set proxy for the validation objective in Eq. 2 and subsumes both an *origin* requirement (the new capability must improve over the previous  $f_C$  on the cases it targets) and a *regression* requirement (it must not degrade currently-correct cases). Because  $\mathcal{D}_{\text{train}}$  and  $\mathcal{D}_{\text{val}}$  are disjoint, this proxy does not contaminate the held-out numbers we report; the only residual concern is selection bias in  $c^*$  over a discrete candidate set, which scales logarithmically in  $M$ , so we keep  $M$  at a single-digit budget. A broader empirical estimate is left to future work. When the action is both, the paired skill additionally serves as the tool’s mastery SOP, gating tool invocation at inference (Section 3.4). The full procedure runs for  $N$  iterations (default  $N=3$ ); pseudocode is given in Appendix A.

### 3.3 Skills

A *skill* is a structured Markdown document with four fields: **When-to-Use** (a trigger predicate over question and image type), **Strategy** (a numbered step-by-step SOP), **Common Pitfalls**, and **Worked Example**. The Generator writes a skill conditioned on the AnalyzerDecider’s root-cause analysis; if a skill covering the same problem class already exists in  $\mathcal{S}$ , new insights are merged into it rather than creating a duplicate, preventing library bloat. At inference, `Retrieve(·)` (Eq. 1) returns the top- $K$  skills by BM25 similarity to the question, which the Solver appends to the system prompt.

### 3.4 Tools and Mastery

A *tool* is a Python function ( $\leq 150$  lines) that takes one or more image paths and returns processed images or extracted data, using standard CV libraries (OpenCV, PIL, NumPy). The Generator writes a tool conditioned on a perceptual bottleneck diagnosis; representative examples produced in our experiments include chart re-rendering with enlarged axis labels, saliency-guided sub-region extraction, and contrast enhancement for low-quality documents.

**Mastery as a paired skill.** When the AnalyzerDecider’s action is both, the Generator emits a skill alongside the tool, and this paired skill is the tool’s mastery SOP. Its When-to-Use predicate specifies the input patterns on which the tool helps, its Common Pitfalls flag the patterns on which it hurts, and its Strategy instructs the Solver how to invoke the tool and consume its output. Because tool invocation is gated by retrieving the paired skill, deployment is selective by construction rather

than indiscriminate—a property that distinguishes DYNAMO from prior tool-creation methods (Qian et al., 2023; Cai et al., 2024) that expose generated tools without learned applicability boundaries.

**External tool sets (Mode B).** When a curated tool set  $\mathcal{T}_0 \neq \emptyset$  is provided, DYNAMO can operate in **Mode B**: skip tool generation and instead synthesise a mastery skill for each provided tool  $t_j \in \mathcal{T}_0$ , learning when to invoke it from the agent’s own behaviour on  $\mathcal{D}_{\text{train}}$ . Mode B requires no code generation, making it practical when tool quality is high but deployment strategy is unknown. Experiment II (Section 4.3) evaluates Mode B.

### 3.5 Online Adaptation to Distribution Shift

In real deployments, DYNAMO runs against a streaming query distribution that can shift over time—a stream may begin with high-resolution perception cases and later start including MathVista-style reasoning cases, or vice versa. We design DYNAMO to detect such shifts and *update* its skills and tools online, without weight updates and without the practitioner having to anticipate the shift or prepare a per-family training set in advance.

**Feedback signal.** Running the evolution loop online requires a per-case correctness signal to drive diagnosis and selection (Eqs. 1–3). Raw production traffic typically lacks ground-truth labels, so DYNAMO assumes one of the practical feedback channels available in real deployments: a small fraction of queries reviewed by humans in the loop, an automated quality-assurance pipeline, or an LLM-based verifier that scores the agent’s answer post-hoc. Our experiments use the benchmark’s ground-truth labels as a stand-in for this channel; the adaptation policy itself is agnostic to the source of the correctness signal, and the framework therefore does not claim fully unsupervised online adaptation.

DYNAMO continuously samples a sub-training set from the most recent queries (together with their correctness signals) and runs the evolution loop of Section 3.2 on them, accumulating skills and tools into the library  $\mathcal{C}$ . Let  $W$  be a rolling-window length. The dominant task family at step  $t$  is

$$\hat{\phi}_t = \arg \max_{\phi} \sum_{i=t-W+1}^t \mathbf{1}[\text{fam}(x_i) = \phi], \quad (4)$$

where  $\text{fam}(\cdot)$  classifies each query from its question and image features. When Eq. 4 flips ( $\hat{\phi}_t \neq \hat{\phi}_{t-1}$ ), the next evolution iteration is triggered on



the new sub-stream and produces updated skills and tools that handle the new distribution. The library is thus kept in sync with the current query distribution. We treat this autonomous online adaptation—decoupled from the source of the correctness signal—as one of the method’s contributions; Experiment I (Section 4.2) evaluates it against a static baseline that never refreshes and an oracle initialized with the full capability set.

## 4 Experiments

### 4.1 Experimental Setup

**Benchmarks.** For the evolution-from-scratch setting (Exp I) we use four visual reasoning benchmarks: **ChartQA** (Masry et al., 2022) (multi-step numerical reasoning over charts; 2,500 test questions), **MathVista** (Lu et al., 2024) (math reasoning in figures and plots; 1,000 testmini questions), **HRBench4K** (Wang et al., 2025) (perception on 4K-resolution images; 800 questions), and **V\*** (Wu and Xie, 2024) (visual search for an object’s property; 191 questions); the first two stress *cognitive* bottlenecks while the latter two stress *perceptual* ones, and all four use accuracy as the metric. For the tool-mastery setting (Exp II) we use **GTA** (Wang et al., 2024b), 229 real-world tool-use tasks across perception, operation, logic, and creativity categories, reporting per-step tool-selection, argument-prediction, and instruction-following accuracy plus end-to-end answer accuracy. For the RL comparison (Exp III) we follow the VTool-R1 protocol (Wu et al., 2025a) on **ChartQA** and **TableQA**, and reuse **V\*** and **HRBench4K** under the DeepEyes protocol (Zheng et al., 2025).

**Foundation models.** Experiment I evaluates five proprietary and open-weight VLMs: **GPT-4o** (?), **o4-mini** (?), **GPT-5.4** (?), **Doubao-Seed-2.0** (?), and **Qwen3.5-27B** (?). Experiment II evaluates **GPT-4o**, **GPT-5.4**, and **Doubao-Seed-2.0** on the GTA benchmark, plus a controlled **Doubao-Seed-2.0-Pro** analysis that inspects how provided-tool skills change with tool abstraction and task environment. Experiment III uses the **Qwen2.5-VL** family (3B, 7B) to align with the VTool-R1 evaluation protocol. All VLM weights are *frozen* throughout; no gradient updates are performed.

**Evolution protocol.** For each benchmark, we sample 10% of the official training split as  $\mathcal{D}_{\text{train}}$  and run the DYNAMO evolution loop for  $N=3$  iterations. The capability set  $\mathcal{C}$  is then frozen and

evaluated on the held-out validation / test split. For each evolution configuration (*Skill Only*, *Tool Only*, *Full*), we run **three independent seeds** that vary the 10% training subset and the Generator’s sampling, and report mean  $\pm$  standard deviation in Table 1. The *None* (Base Agent) row runs the frozen VLM with no evolved capabilities once per cell, so no standard deviation is reported.

**Baselines.** For Experiment I, we compare four evolution configurations that isolate the contribution of each component: **None (Base Agent)** invokes the frozen VLM with no evolved capabilities (the zero-shot baseline); **Skill Only** restricts the Generator to producing reasoning skills (no tool generation); **Tool Only** restricts the Generator to producing executable visual tools (no skill generation); **Full** is the unrestricted DYNAMO system, where the AnalyzerDecider chooses on each case whether to generate a skill, a tool, or both.

**Evaluation metrics.** All primary metrics are task-level accuracy on the held-out split. For Experiment II we report four of GTA’s five canonical metrics: three step-by-step metrics—**InstAcc** (Instruction-Following Accuracy: fraction of steps executed without errors), **ToolAcc** (Tool Selection Accuracy), and **ArgAcc** (Argument Prediction Accuracy)—and the end-to-end **AnsAcc** (Answer Accuracy on full tool-using execution). For the supplementary controlled GTA analysis, we also report tool usage rate, average number of tool calls, and dominant tool-call chains.

### 4.2 Experiment I: Autonomous Evolution from Scratch

We test whether DYNAMO can bootstrap a useful capability library from scratch ( $\mathcal{S}=\emptyset$ ,  $\mathcal{T}=\emptyset$ ) across five backbones and four benchmarks, and whether the library remains useful when the query distribution shifts over time (Section 3.5).

**Per-benchmark results.** Table 1 reports accuracy for the four configurations on each backbone. Comparing the Full system to the zero-shot *None* baseline, DYNAMO improves direct inference on all 20 model–benchmark cells, with an average gain of +5.6 accuracy points. The largest gains are on o4-mini / V\* (+14.1), o4-mini / HRBench4K (+12.7), GPT-5.4 / HRBench4K (+10.7), Doubao-Seed-2.0 / V\* (+8.8), and Qwen3.5-27B / HRBench4K (+8.4). The relative strength of *Skill Only* and *Tool Only* also follows the cognitive / perceptual split:

| Model           | Evolution Strategy | HRBench4K     | MathVista     | V*            | ChartQA       |
|-----------------|--------------------|---------------|---------------|---------------|---------------|
| GPT-4o          | None (Base Agent)  | 0.6386        | 0.6711        | 0.5828        | 0.7552        |
|                 | Skill Only         | 0.6733±0.0121 | 0.7122±0.0038 | 0.6424±0.0066 | 0.7785±0.0026 |
|                 | Tool Only          | 0.6452±0.0044 | 0.6900±0.0029 | 0.6225±0.0066 | 0.7722±0.0029 |
|                 | Full               | 0.6886±0.0071 | 0.7189±0.0029 | 0.6689±0.0239 | 0.7799±0.0045 |
| o4-mini         | None (Base Agent)  | 0.73          | 0.753333      | 0.72847       | 0.7833        |
|                 | Skill Only         | 0.7724±0.0058 | 0.7470±0.0107 | 0.7616±0.0066 | 0.7866±0.0151 |
|                 | Tool Only          | 0.7457±0.1670 | 0.7356±0.0328 | 0.8387±0.0629 | 0.7920±0.0213 |
|                 | Full               | 0.8571±0.0099 | 0.7559±0.0055 | 0.8698±0.0076 | 0.8007±0.0081 |
| GPT-5.4         | None (Base Agent)  | 0.7471        | 0.7589        | 0.7748        | 0.7974        |
|                 | Skill Only         | 0.7805±0.0036 | 0.7752±0.0039 | 0.7682±0.0199 | 0.8090±0.0022 |
|                 | Tool Only          | 0.8367±0.0346 | 0.7663±0.0061 | 0.8477±0.0132 | 0.8108±0.0052 |
|                 | Full               | 0.8538±0.0036 | 0.7815±0.0101 | 0.8521±0.0233 | 0.8276±0.0278 |
| Doubao-Seed-2.0 | None (Base Agent)  | 0.8214        | 0.8544        | 0.8675        | 0.8152        |
|                 | Skill Only         | 0.8600±0.0049 | 0.8556±0.0109 | 0.8609±0.0115 | 0.8198±0.0031 |
|                 | Tool Only          | 0.9005±0.0022 | 0.8659±0.0061 | 0.9448±0.0038 | 0.8156±0.0068 |
|                 | Full               | 0.8990±0.0064 | 0.8681±0.0105 | 0.9558±0.0038 | 0.8347±0.0295 |
| Qwen3.5-27B     | None (Base Agent)  | 0.8171        | 0.8088        | 0.88079       | 0.8114        |
|                 | Skill Only         | 0.8721±0.0030 | 0.8152±0.0080 | 0.8675±0.0066 | 0.8097±0.0090 |
|                 | Tool Only          | 0.8986±0.0000 | 0.8241±0.0034 | 0.9514±0.0101 | 0.8115±0.0041 |
|                 | Full               | 0.9010±0.0008 | 0.8252±0.0023 | 0.9603±0.0175 | 0.8158±0.0026 |

Table 1: **Autonomous evolution from scratch.** Accuracy across five VLM backbones and four benchmarks. The Full DYNAMO configuration improves over the zero-shot *None* baseline on all 20 model–benchmark cells.

Skill Only is competitive on ChartQA and MathVista, where the gain comes from better decomposition and calculation procedures, while Tool Only and Full dominate on HRBench4K and V\*, where sub-region inspection and visual search are central. The only cell where an ablation edges out Full is Doubao / HRBench4K, where Tool Only beats Full by 0.0015—an expected consequence of leaving capability choice to the AnalyzerDecider rather than enforcing both per case.

**Capability quality.** Appendix D inspects three representative forms of evolved capability: a ReFocus-style chart/table editing case, a ZoomEye-style coarse-to-fine search case, and an interleaved skill that pairs a textual SOP with visual evidence. The case studies illustrate what DYNAMO generates; they do not claim to reproduce the full algorithms of ReFocus or ZoomEye.

**Online adaptation to distribution shift.** To simulate real online user traffic—where the task distribution grows and shifts over time as different users and requests arrive—we replay a 208-step non-stationary stream in which each phase introduces a new task family while keeping carry-over cases from earlier phases: starting with HRBench4K, then adding MathVista, ChartQA, and finally V\*

as the new dominant family. The replay reuses completed per-case outputs, and the benchmark’s ground-truth labels stand in for the per-case correctness signal that a real deployment would obtain from human review, automated QA, or an LLM verifier (Section 3.5). This isolates the adaptation policy from the cost of regenerating capabilities online and from the cost of the feedback channel itself; we do not claim fully unsupervised online adaptation. We compare *Direct* (no capabilities, same as the *None* baseline in Table 1), *Static* (capabilities evolved on the first phase only and never refreshed), *Online adapt* (DYNAMO detects the shift via the rolling-window rule of Section 3.5 and re-runs the evolution loop on the new sub-stream), and *Oracle* (pre-evolved on every family; upper bound).

Figure 3 compares the four policies across all five backbones and three stream constructions: the *natural* mixture (a fixed-seed mix of the four families), *capability-relevant* (cases whose outcome depends on the matching capability), and *stress* (cases the agent fails on without the matching capability). The ordering  $Direct \leq Static < Online\ adapt \approx Oracle$  is consistent across every backbone  $\times$  protocol cell. On GPT-5.4 the gap is representative: *Online adapt* reaches 0.83 vs. 0.74 static on natural; widens to 0.89 vs. 0.65 on capability-relevant;




Figure 3: **Online adaptation to distribution shift.** Replay accuracy across five backbones and three stream constructions. *Online adapt* stays close to the *Oracle* upper bound and far above *Static* everywhere; the gap is largest under the stress protocol, where *Direct* nearly collapses. Error bars are seed std (50–200 seeds per cell).

and on stress, *Direct* collapses to 0.03 while *Online adapt* still recovers to 0.95. The per-step trajectory along the stream, with shift-detection latencies of 2–7 cases per phase, is shown in Appendix C.2.

### 4.3 Experiment II: Learning to Use Pre-Provided Tools

**Results on GTA.** Table 2 compares three methods on the 121-case GTA validation split: *Base Agent* (zero-shot VLM with the GTA tools exposed); an *XSkill-style* baseline (Jiang et al., 2026) that adopts XSkill’s prompt-injection mechanism on top of the same tool environment (see Appendix A.1 for the port); and DYNAMO Mode B (skills + mastery learned on the training set). DYNAMO improves every metric on every backbone, with the largest gains on Doubao-Seed-2.0 (+18.7 ToolAcc, +15.8 ArgAcc, +16.1 InstAcc, +4.9 AnsAcc) and smaller but uniformly positive gains on GPT-4o (+4.3 / +2.2 / +11.4 / +5.0) and GPT-5.4 (+6.8 / +3.2 / +6.1 / +1.7).

**What does environment-specific mastery contribute?** The *XSkill-style* baseline isolates generic skill injection from environment-specific mastery: it uses the same tools and scorer but its skill is a generic visual-reasoning prompt without GTA’s tool-chain, argument, or answer-normalisation protocols. On every backbone, DYNAMO beats XSkill-style by 11–20 points on each step-level metric and 2.5–4.9 on AnsAcc. More tellingly, *XSkill-style* itself drops below *Base* on the step-level metrics for GPT-4o and GPT-5.4—adding a generic visual-reasoning prompt without learning the environment’s tool protocol can actually confuse tool selection. The gain over

| Model           | Method        | ToolAcc     | ArgAcc      | InstAcc     | AnsAcc      |
|-----------------|---------------|-------------|-------------|-------------|-------------|
| GPT-4o          | Base Agent    | 80.3        | 62.0        | 52.0        | 66.1        |
|                 | XSkill-style  | 71.3        | 53.0        | 49.5        | 68.6        |
|                 | DYNAMO (Ours) | <b>84.6</b> | <b>64.2</b> | <b>63.4</b> | <b>71.1</b> |
| GPT-5.4         | Base Agent    | 79.6        | 60.2        | 57.0        | 72.7        |
|                 | XSkill-style  | 66.3        | 48.0        | 46.2        | 71.1        |
|                 | DYNAMO (Ours) | <b>86.4</b> | <b>63.4</b> | <b>63.1</b> | <b>74.4</b> |
| Doubao-Seed-2.0 | Base Agent    | 67.7        | 45.5        | 41.6        | 76.9        |
|                 | XSkill-style  | 72.8        | 48.4        | 44.4        | 76.9        |
|                 | DYNAMO (Ours) | <b>86.4</b> | <b>61.3</b> | <b>57.7</b> | <b>81.8</b> |

Table 2: **GTA results: Base Agent, XSkill-style, and DYNAMO Mode B on the 121-case validation split.** Accuracy (%) per backbone; best per row in bold.

both baselines is therefore attributable not to skill injection per se, but to mastery skills that encode the target environment’s specific tool chains, argument formats, and answer normalisation rules.

**Controlled tool-policy analysis.** Table 3 runs a complementary controlled study on GTA with Doubao-Seed-2.0-Pro to inspect what the learned policy actually does when either the tool abstraction or the task environment changes. Two findings stand out. First, fixing the environment to OCR+Calculator, stronger composite tools compress the policy: atomic tools reach 86.7 via an OCR→Calculator chain (22/30 cases), while VisualArithmeticSolver reaches 90.0 with a one-call policy (30/30). Second, the same OCR tool plays different roles in different environments: in OCR+Calc it feeds arithmetic (OCR→Calculator, 19/20), while in OCR+Search it feeds external lookup (OCR→Search→OCR→Search, 20/20).

| Study         | Condition        | Cases | Acc. | Tool use | Avg. calls |
|---------------|------------------|-------|------|----------|------------|
| Tool strength | Direct, no tools | 30    | 83.3 | 0.0      | 0.00       |
| Tool strength | Atomic OCR+Calc  | 30    | 86.7 | 100.0    | 1.83       |
| Tool strength | Composite VAS    | 30    | 90.0 | 100.0    | 1.00       |
| OCR role      | OCR+Calc env     | 20    | 80.0 | 100.0    | 2.20       |
| OCR role      | OCR+Search env   | 20    | 85.0 | 100.0    | 2.35       |

Table 3: **Controlled tool-policy analysis on GTA.** Acc. and Tool use in %; Avg. calls is a count.

#### 4.4 Experiment III: RL and Self-Evolution Across Visual Tool Tasks

We compare DYNAMO against task-specific RL on two task families using two Qwen2.5-VL backbones (3B, 7B): structured visual QA (ChartQA and TableQA) under the VTool-R1 protocol (Wu et al., 2025a), and high-resolution perception (V\* and HRBench4K) under the DeepEyes protocol (Zheng et al., 2025).

**ChartQA and TableQA.** Using the gap-recovery fraction  $(\text{DYNAMO} - \text{Direct})/(\text{RL} - \text{Direct})$ , DYNAMO recovers 65.7% and 99.4% of the ChartQA gap at 3B and 7B respectively, and 67.5% and 91.4% of the TableQA gap. DYNAMO + RL further beats RL alone by 1–3 points on every cell, showing the two interventions compose additively.

**High-resolution visual search.** On V\*, DYNAMO alone beats task-specific RL at both scales (+6.3 at 3B, +1.0 at 7B). On HRBench4K, DYNAMO beats RL at 3B by +4.0 but trails at 7B by 1.8. Adding DYNAMO on top of the RL-trained policy raises 7B V\* to 92.7 and 7B HRBench4K to 81.3, the strongest configurations in those cells.

**Cost.** The RL baselines need on the order of  $10^4$ – $10^5$  rollouts and 80–240 gradient updates per backbone (Appendix A.2 for the per-protocol breakdown). DYNAMO performs no weight updates: a run uses 10% of the training split for  $N=3$  passes of frozen-VLM API calls, a few thousand calls per benchmark (token totals in Appendix A).

**Takeaway.** DYNAMO’s self-evolved skills and tools alone reach RL-level accuracy on both task families (most of the gap on ChartQA/TableQA; matching or beating RL on V\* and at 3B on HRBench4K) without any weight updates. RL and self-evolution are not at odds: DYNAMO + RL is the strongest configuration in every cell, so the two interventions compose additively.

| Backbone      | Variant     | ChartQA | TableQA | V*    | HRBench4K |
|---------------|-------------|---------|---------|-------|-----------|
| Qwen2.5-VL-3B | Direct      | 45.40   | 43.09   | 71.20 | 63.62     |
|               | RL          | 66.95   | 56.25   | 78.01 | 66.50     |
|               | DYNAMO      | 59.56   | 51.97   | 84.29 | 70.50     |
|               | DYNAMO + RL | 68.28   | 59.54   | 85.34 | 71.25     |
| Qwen2.5-VL-7B | Direct      | 56.05   | 56.91   | 79.06 | 70.50     |
|               | RL          | 75.18   | 68.42   | 84.82 | 76.88     |
|               | DYNAMO      | 75.060  | 67.43   | 85.86 | 75.12     |
|               | DYNAMO + RL | 76.88   | 69.41   | 92.67 | 81.25     |

Table 4: **Self-evolution vs. task-specific RL across two task families.** Acc. (%) on Qwen2.5-VL 3B and 7B.

#### 4.5 Ablation Studies

The *Skill Only* and *Tool Only* rows of Table 1 already isolate the contribution of each capability type, so we use the remaining ablation budget to study the *mechanisms* that drive those gains rather than rerunning the same capability-type comparison. Appendix B reports three artifact-backed mechanism probes that support the claims we defend with the current budget: (i) generic prompt injection does not explain the HRBench4K gain—only the validated visual-tool path does; (ii) text-only diagnosis cannot materialise a usable image-processing tool; and (iii) diagnosing only correct or only incorrect ChartQA cases fails to drive useful evolution, supporting the full loop’s use of contrast between solved and unsolved cases. The remaining mechanism removal (multi-candidate exploration) needs a larger subset and is left to future work.

## 5 Conclusion

DYNAMO is a training-free framework in which a frozen VLM agent evolves a task-family-specific library of reasoning skills and executable visual tools from a small labeled training subset, with the Analyzer choosing on each case whether to emit a skill, a tool, or both. A mastery skill paired with each generated tool gates its invocation at inference, and an online adaptation loop refreshes the library when the query distribution shifts.

Across four visual-reasoning benchmarks and five backbones, DYNAMO improves direct inference on all 20 model–benchmark cells; on GTA it lifts every step-level tool-use metric across all tested backbones; and against task-specific RL (VTool-R1, DeepEyes), DYNAMO matches or beats RL on high-resolution perception, recovers most of the structured-VQA gap, and composes additively with RL when both are available. These results suggest much of task-specific VLM adaptation can be obtained without gradient updates, by letting the agent shape its capability library.



## Limitations

**Diagnosis bounded by the backbone.** The ANALYZERDECIDER’s bottleneck classification (*cognitive* vs. *perceptual*) and the GENERATOR’s candidates are produced by the same frozen VLM that the agent uses. On backbones with weak self-introspection a mis-classified bottleneck can propagate to a mis-targeted skill or tool; the promotion guard rejects under-performing candidates but does not correct upstream diagnosis errors, so smaller or less capable backbones may converge more slowly or to a less useful library.

**Benchmark scope.** We evaluate capability evolution on image-based VLMs across four visual-reasoning benchmarks plus GTA and the VTool-R1 / DeepEyes splits. Video VLMs and other modalities (3D spatial reasoning, medical imaging) are outside the current scope; whether the cognitive / perceptual decision rule and the executable-tool interface transfer to those settings is left to future work.

**Reliance on a per-case correctness signal.** Both the offline evolution loop and the online adaptation policy (Sections 3.2–3.5) require a per-case correctness signal to drive diagnosis and the multi-candidate selection rule. Our experiments use ground-truth labels from the benchmark splits as this signal, which corresponds to deployment settings with human-in-the-loop review, an automated quality-assurance pipeline, or an LLM-based post-hoc verifier. Deployments that have no access to any of these channels—i.e. fully unsupervised production traffic—fall outside the scope of the current method; extending DYNAMO to that setting would require a confidence-based or self-consistency-based feedback substitute that we leave to future work.

## References

- Tianle Cai, Xuezhi Wang, Tengyu Ma, Xinyun Chen, and Denny Zhou. 2024. Large language models as tool makers. In *ICLR*. OpenReview.net.
- Ruoxi Cheng, Haoxuan Ma, Teng Ma, and Hongyi Zhang. 2025. Ecoalign: An economically rational framework for efficient LVLM alignment. *CoRR*, abs/2511.11301.
- Xingyu Fu, Minqian Liu, Zhengyuan Yang, John Corring, Yijuan Lu, Jianwei Yang, Dan Roth, Dinei A. F. Florêncio, and Cha Zhang. 2025. Refocus: Visual editing as a chain of thought for structured image

understanding. In *ICML*, Proceedings of Machine Learning Research. PMLR / OpenReview.net.

- Tanmay Gupta and Aniruddha Kembhavi. 2023. Visual programming: Compositional visual reasoning without training. In *CVPR*, pages 14953–14962. IEEE.

Jack Hong, Chenxiao Zhao, ChengLin Zhu, Weiheng Lu, Guohai Xu, and Xing Yu. 2025. *Deep-eyesv2: Toward agentic multimodal model*. *CoRR*, abs/2511.05271.

Guanyu Jiang, Zhaochen Su, Xiaoye Qu, and Yi R. Fung. 2026. Xskill: Continual learning from experience and skills in multimodal agents. *CoRR*, abs/2603.12056.

Zongxia Li, Xiyang Wu, Hongyang Du, Fuxiao Liu, Huy Nghiem, and Guangyao Shi. 2025. *A survey of state of the art large vision language models: Benchmark evaluations and challenges*. In *IEEE/CVF Conference on Computer Vision and Pattern Recognition Workshops, CVPR Workshops 2025, Nashville, TN, USA, June 11-15, 2025*, pages 1587–1606. Computer Vision Foundation / IEEE.

Pan Lu, Hritik Bansal, Tony Xia, Jiacheng Liu, Chunyuan Li, Hannaneh Hajishirzi, Hao Cheng, Kai-Wei Chang, Michel Galley, and Jianfeng Gao. 2024. Mathvista: Evaluating mathematical reasoning of foundation models in visual contexts. In *ICLR*. OpenReview.net.

Pan Lu, Baolin Peng, Hao Cheng, Michel Galley, Kai-Wei Chang, Ying Nian Wu, Song-Chun Zhu, and Jianfeng Gao. 2023. Chameleon: Plug-and-play compositional reasoning with large language models. In *NeurIPS*.

Haoxuan Ma, Guannan Lai, and Han-Jia Ye. 2026. Mmr-bench: A comprehensive benchmark for multimodal LLM routing. *CoRR*, abs/2601.17814.

Ahmed Masry, Do Xuan Long, Jia Qing Tan, Shafiq R. Joty, and Enamul Hoque. 2022. Chartqa: A benchmark for question answering about charts with visual and logical reasoning. In *ACL (Findings)*, Findings of ACL, pages 2263–2279. Association for Computational Linguistics.

Jingwei Ni, Yihao Liu, Xinpeng Liu, Yutao Sun, Mengyu Zhou, Pengyu Cheng, Dexin Wang, Erchao Zhao, Xiaoxi Jiang, and Guanjun Jiang. 2026. Trace2skill: Distill trajectory-local lessons into transferable agent skills. *CoRR*, abs/2603.25158.

Cheng Qian, Chi Han, Yi Ren Fung, Yujia Qin, Zhiyuan Liu, and Heng Ji. 2023. CREATOR: tool creation for disentangling abstract and concrete reasoning of large language models. In *EMNLP (Findings)*, Findings of ACL, pages 6922–6939. Association for Computational Linguistics.

Runqi Qiao, Qiuna Tan, Minghan Yang, Guanting Dong, Peiqing Yang, Shiqiang Lang, Enhui Wan, Xiaowan Wang, Yida Xu, Lan Yang, Chong Sun, Chen Li,

|     |                                                                                                                                                                                                                                                                                                                                                           |     |
|-----|-----------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|-----|
| 688 | and Honggang Zhang. 2025. V-thinker: Interactive thinking with images. <i>CoRR</i> , abs/2511.04460.                                                                                                                                                                                                                                                      |     |
| 689 |                                                                                                                                                                                                                                                                                                                                                           |     |
| 690 | Yujia Qin, Shihao Liang, Yining Ye, Kunlun Zhu, Lan Yan, Yaxi Lu, Yankai Lin, Xin Cong, Xiangru Tang, Bill Qian, Sihan Zhao, Lauren Hong, Runchu Tian, Ruobing Xie, Jie Zhou, Mark Gerstein, Dahai Li, Zhiyuan Liu, and Maosong Sun. 2024. Toolllm: Facilitating large language models to master 16000+ real-world apis. In <i>ICLR</i> . OpenReview.net. |     |
| 691 |                                                                                                                                                                                                                                                                                                                                                           |     |
| 692 |                                                                                                                                                                                                                                                                                                                                                           |     |
| 693 |                                                                                                                                                                                                                                                                                                                                                           |     |
| 694 |                                                                                                                                                                                                                                                                                                                                                           |     |
| 695 |                                                                                                                                                                                                                                                                                                                                                           |     |
| 696 |                                                                                                                                                                                                                                                                                                                                                           |     |
| 697 | Haozhan Shen, Kangjia Zhao, Tiancheng Zhao, Ruochen Xu, Zilun Zhang, Mingwei Zhu, and Jianwei Yin. 2025. Zoomeye: Enhancing multimodal llms with human-like zooming capabilities through tree-based image exploration. In <i>EMNLP</i> , pages 6602–6618. Association for Computational Linguistics.                                                      |     |
| 698 |                                                                                                                                                                                                                                                                                                                                                           |     |
| 699 |                                                                                                                                                                                                                                                                                                                                                           |     |
| 700 |                                                                                                                                                                                                                                                                                                                                                           |     |
| 701 |                                                                                                                                                                                                                                                                                                                                                           |     |
| 702 |                                                                                                                                                                                                                                                                                                                                                           |     |
| 703 |                                                                                                                                                                                                                                                                                                                                                           |     |
| 704 | Noah Shinn, Federico Cassano, Ashwin Gopinath, Karthik Narasimhan, and Shunyu Yao. 2023. Reflexion: language agents with verbal reinforcement learning. In <i>NeurIPS</i> .                                                                                                                                                                               |     |
| 705 |                                                                                                                                                                                                                                                                                                                                                           |     |
| 706 |                                                                                                                                                                                                                                                                                                                                                           |     |
| 707 |                                                                                                                                                                                                                                                                                                                                                           |     |
| 708 | Alex Su, Haozhe Wang, Weiming Ren, Fangzhen Lin, and Wenhui Chen. 2025. Pixel reasoner: Incentivizing pixel-space reasoning with curiosity-driven reinforcement learning. <i>CoRR</i> , abs/2505.15966.                                                                                                                                                   |     |
| 709 |                                                                                                                                                                                                                                                                                                                                                           |     |
| 710 |                                                                                                                                                                                                                                                                                                                                                           |     |
| 711 |                                                                                                                                                                                                                                                                                                                                                           |     |
| 712 | Yutao Sun, Mingshuai Chen, Tiancheng Zhao, Ruochen Xu, Zilun Zhang, and Jianwei Yin. 2025. The self-improvement paradox: Can language models bootstrap reasoning capabilities without external scaffolding? In <i>ACL (Findings)</i> , Findings of ACL, pages 6501–6512. Association for Computational Linguistics.                                       |     |
| 713 |                                                                                                                                                                                                                                                                                                                                                           |     |
| 714 |                                                                                                                                                                                                                                                                                                                                                           |     |
| 715 |                                                                                                                                                                                                                                                                                                                                                           |     |
| 716 |                                                                                                                                                                                                                                                                                                                                                           |     |
| 717 |                                                                                                                                                                                                                                                                                                                                                           |     |
| 718 |                                                                                                                                                                                                                                                                                                                                                           |     |
| 719 | Dídac Surís, Sachit Menon, and Carl Vondrick. 2023. Vipergpt: Visual inference via python execution for reasoning. In <i>ICCV</i> , pages 11854–11864. IEEE.                                                                                                                                                                                              |     |
| 720 |                                                                                                                                                                                                                                                                                                                                                           |     |
| 721 |                                                                                                                                                                                                                                                                                                                                                           |     |
| 722 | Guanzhi Wang, Yuqi Xie, Yunfan Jiang, Ajay Mandlekar, Chaowei Xiao, Yuke Zhu, Linxi Fan, and Anima Anandkumar. 2024a. Voyager: An open-ended embodied agent with large language models. <i>Trans. Mach. Learn. Res.</i> , 2024.                                                                                                                           |     |
| 723 |                                                                                                                                                                                                                                                                                                                                                           |     |
| 724 |                                                                                                                                                                                                                                                                                                                                                           |     |
| 725 |                                                                                                                                                                                                                                                                                                                                                           |     |
| 726 |                                                                                                                                                                                                                                                                                                                                                           |     |
| 727 | Jize Wang, Zerun Ma, Yining Li, Songyang Zhang, Cailian Chen, Kai Chen, and Xinyi Le. 2024b. GTA: A benchmark for general tool agents. In <i>NeurIPS</i> .                                                                                                                                                                                                |     |
| 728 |                                                                                                                                                                                                                                                                                                                                                           |     |
| 729 |                                                                                                                                                                                                                                                                                                                                                           |     |
| 730 | Wenbin Wang, Liang Ding, Minyan Zeng, Xiabin Zhou, Li Shen, Yong Luo, Wei Yu, and Dacheng Tao. 2025. Divide, conquer and combine: A training-free framework for high-resolution image perception in multimodal large language models. In <i>AAAI</i> , pages 7907–7915. AAAI Press.                                                                       |     |
| 731 |                                                                                                                                                                                                                                                                                                                                                           |     |
| 732 |                                                                                                                                                                                                                                                                                                                                                           |     |
| 733 |                                                                                                                                                                                                                                                                                                                                                           |     |
| 734 |                                                                                                                                                                                                                                                                                                                                                           |     |
| 735 |                                                                                                                                                                                                                                                                                                                                                           |     |
| 736 | Mingyuan Wu, Jingcheng Yang, Jize Jiang, Meitang Li, Kaizhuo Yan, Hanchao Yu, Minjia Zhang, Chengxiang Zhai, and Klara Nahrstedt. 2025a. Vtool-r1: Vlms learn to think with images via reinforcement learning on multimodal tool use. <i>CoRR</i> , abs/2505.19255.                                                                                       |     |
| 737 |                                                                                                                                                                                                                                                                                                                                                           |     |
| 738 |                                                                                                                                                                                                                                                                                                                                                           |     |
| 739 |                                                                                                                                                                                                                                                                                                                                                           |     |
| 740 |                                                                                                                                                                                                                                                                                                                                                           |     |
| 741 |                                                                                                                                                                                                                                                                                                                                                           |     |
|     | Penghao Wu and Saining Xie. 2024. V*: Guided visual search as a core mechanism in multimodal llms. In <i>CVPR</i> , pages 13084–13094. IEEE.                                                                                                                                                                                                              | 742 |
|     |                                                                                                                                                                                                                                                                                                                                                           | 743 |
|     |                                                                                                                                                                                                                                                                                                                                                           | 744 |
|     | Rong Wu, Xiaoman Wang, Jianbiao Mei, Pinlong Cai, Daocheng Fu, Cheng Yang, Licheng Wen, Xueming Yang, Yufan Shen, Yuxin Wang, and Botian Shi. 2025b. Evolver: Self-evolving LLM agents through an experience-driven lifecycle. <i>CoRR</i> , abs/2510.16079.                                                                                              | 745 |
|     |                                                                                                                                                                                                                                                                                                                                                           | 746 |
|     |                                                                                                                                                                                                                                                                                                                                                           | 747 |
|     |                                                                                                                                                                                                                                                                                                                                                           | 748 |
|     |                                                                                                                                                                                                                                                                                                                                                           | 749 |
|     |                                                                                                                                                                                                                                                                                                                                                           | 750 |
|     | Zhangchen Xu, Fengqing Jiang, Luyao Niu, Yuntian Deng, Radha Poovendran, Yejin Choi, and Bill Yuchen Lin. 2025. Magpie: Alignment data synthesis from scratch by prompting aligned llms with nothing. In <i>The Thirteenth International Conference on Learning Representations, ICLR 2025, Singapore, April 24-28, 2025</i> . OpenReview.net.            | 751 |
|     |                                                                                                                                                                                                                                                                                                                                                           | 752 |
|     |                                                                                                                                                                                                                                                                                                                                                           | 753 |
|     |                                                                                                                                                                                                                                                                                                                                                           | 754 |
|     |                                                                                                                                                                                                                                                                                                                                                           | 755 |
|     |                                                                                                                                                                                                                                                                                                                                                           | 756 |
|     |                                                                                                                                                                                                                                                                                                                                                           | 757 |
|     | Xu Yang, Yingzhe Peng, Haoxuan Ma, Shuo Xu, Chi Zhang, Yucheng Han, and Hanwang Zhang. 2024. Lever LM: configuring in-context sequence to lever large vision language models. In <i>NeurIPS</i> .                                                                                                                                                         | 758 |
|     |                                                                                                                                                                                                                                                                                                                                                           | 759 |
|     |                                                                                                                                                                                                                                                                                                                                                           | 760 |
|     |                                                                                                                                                                                                                                                                                                                                                           | 761 |
|     | Ruixiang Zhang, Richard He Bai, Huangjie Zheng, Navdeep Jaitly, Ronan Collobert, and Yizhe Zhang. 2026. Embarrassingly simple self-distillation improves code generation. <i>CoRR</i> , abs/2604.01193.                                                                                                                                                   | 762 |
|     |                                                                                                                                                                                                                                                                                                                                                           | 763 |
|     |                                                                                                                                                                                                                                                                                                                                                           | 764 |
|     |                                                                                                                                                                                                                                                                                                                                                           | 765 |
|     | Andrew Zhao, Daniel Huang, Quentin Xu, Matthieu Lin, Yong-Jin Liu, and Gao Huang. 2024. Expel: LLM agents are experiential learners. In <i>AAAI</i> , pages 19632–19642. AAAI Press.                                                                                                                                                                      | 766 |
|     |                                                                                                                                                                                                                                                                                                                                                           | 767 |
|     |                                                                                                                                                                                                                                                                                                                                                           | 768 |
|     |                                                                                                                                                                                                                                                                                                                                                           | 769 |
|     | Ziwei Zheng, Michael Yang, Jack Hong, Chenxiao Zhao, Guohai Xu, Le Yang, Chao Shen, and Xing Yu. 2025. Deepeyes: Incentivizing "thinking with images" via reinforcement learning. <i>CoRR</i> , abs/2505.14362.                                                                                                                                           | 770 |
|     |                                                                                                                                                                                                                                                                                                                                                           | 771 |
|     |                                                                                                                                                                                                                                                                                                                                                           | 772 |
|     |                                                                                                                                                                                                                                                                                                                                                           | 773 |
|     |                                                                                                                                                                                                                                                                                                                                                           | 774 |

## A Implementation Details

**Algorithm.** Algorithm 1 gives the full DYNAMO evolution procedure described in Section 3.2.

---

### Algorithm 1: DYNAMO evolution loop.

---

**Input:**  $\mathcal{D}_{\text{train}}$ , frozen VLM  $\pi_\theta$ , iterations  $N$ , candidates  $M$

**Output:** Capability set  $\mathcal{C} = (\mathcal{S}, \mathcal{T})$

```

1  $\mathcal{C} \leftarrow (\emptyset, \emptyset)$ ;
2 for  $n = 1$  to  $N$  do
3   Solve  $\mathcal{D}_{\text{train}}$  with  $f_{\mathcal{C}}$ ; collect attempts;
4    $\mathcal{D}_{\text{sub}} \leftarrow$  sample from  $\mathcal{D}_{\text{train}}$ ;
5    $a, r \leftarrow$ 
     ANALYZERDECIDER( $\mathcal{D}_{\text{sub}}$ ;  $\pi_\theta$ );
6    $\{(s_m, t_m)\}_{m=1}^M \leftarrow$ 
     GENERATOR( $a, r$ ;  $\pi_\theta$ ); //  $s_m$ 
     paired with  $t_m$  acts as its
     mastery SOP
7    $m^* \leftarrow$ 
      $\arg \max_m \text{Acc}(f_{\mathcal{C} \cup (s_m, t_m)}; \mathcal{D}_{\text{train}})$ ;
8   if accuracy improves then
9      $\mathcal{C} \leftarrow \mathcal{C} \cup \{(s_{m^*}, t_{m^*})\}$ ;
10  end
11 end
12 return  $\mathcal{C}$ ;
```

---

**Role prompts.** DYNAMO uses three role-specific system prompts for the same base VLM: *Solver* (standard QA prompt with capability context injected), *AnalyzerDecider* (root-cause analysis and decision), and *Generator* (skill or tool generation). All prompts are provided in full at [\[anonymised repository\]](#).

**Skill retrieval.** At inference time, the Solver retrieves the top- $K=3$  skills from  $\mathcal{S}$  using BM25 cosine similarity over skill titles and When-to-Use fields. Retrieved skill text is appended to the system prompt.

**Tool invocation.** The Solver inspects each tool’s mastery SOP to decide whether to invoke it on the current input. If the input matches a tool’s supported patterns, the tool is applied and its output image(s) are appended to the multimodal context. Tool execution is sandboxed with a 30-second timeout and resource limits.

**Hyperparameters.** Table 5 lists all hyperparameter settings.

Table 5: Hyperparameter settings.

| Parameter         | Value              | Description                                                 |
|-------------------|--------------------|-------------------------------------------------------------|
| Training subset   | 10% of train split | Per-benchmark evolution subset $\mathcal{D}_{\text{train}}$ |
| $N$               | 3                  | Evolution iterations                                        |
| $K$               | 3                  | Skills retrieved per query                                  |
| Max tool lines    | 150                | Max lines for generated Python tool                         |
| Temperature       | 0.7                | Generator temperature                                       |
| Max tokens (gen.) | 2,048              | Max tokens for capability generation                        |

**Compute.** Experiments use the frozen VLM backbones listed in Section 4.1. Each training case incurs at most three frozen-VLM calls (Solver, AnalyzerDecider, and Generator), so the per-benchmark call budget scales linearly in the size of the training subset and in  $N$ . No GPU training is performed.

### A.1 XSkill-style Baseline

**What XSkill is.** XSkill (Jiang et al., 2026) is a continual-learning framework that maintains a library of structured Markdown skills (Definition 2.1 of their paper: each skill carries metadata, a workflow, code templates, and watchpoints) and a complementary library of short JSON experiences (Definition 2.2:  $(c, a, v_e)$  tuples of triggering condition, recommended action, and semantic embedding,  $\leq 64$  words combined). At inference, XSkill decomposes the task into subtasks, retrieves top- $k$  experiences by embedding similarity to the current visual context, rewrites them, prunes irrelevant sections of a retrieved skill against the current image, and injects the adapted skill plus rewritten experiences into the agent’s system prompt as a non-prescriptive reference.

**Our port to GTA.** Our XSkill-style baseline keeps XSkill’s prompt-injection point in the system prompt, the same GTA tool environment, and the same scorer as DYNAMO Mode B. It does not run XSkill’s continual-learning loop end-to-end on GTA training data; instead, the slot that XSkill would fill with a learned Markdown skill is filled by a single generic visual-reasoning skill prompt (no GTA-specific tool-chain, argument-format, or answer-normalisation protocol). This isolates the contribution of XSkill’s skill-injection mechanism from DYNAMO’s GTA-specific mastery skills.

## A.2 RL Baseline Training Compute

Each RL baseline in Experiment III is the published, pre-trained checkpoint from its original protocol, evaluated at both Qwen2.5-VL scales (3B, 7B).

**VTool-R1 (ChartQA, TableQA).** The released checkpoints are reported after 80 GRPO steps with a batch of 32 prompts and  $n_{\text{rollout}}=8$  rollouts per prompt, giving approximately  $80 \times 32 \times 8 = 20,480$  trajectories and 80 gradient updates per backbone.

**DeepEyes (V\*, HRBench4K).** The released checkpoints are reported after 240 RL steps with a batch of 128 prompts and  $n_{\text{rollout}}=4$  rollouts per prompt, giving approximately  $240 \times 128 \times 4 = 122,880$  trajectories and 240 gradient updates per backbone.

**DYNAMO.** A DYNAMO run performs no gradient updates. It uses 10% of each benchmark’s training split as  $\mathcal{D}_{\text{train}}$  and runs the evolution loop for  $N=3$  passes, invoking the frozen VLM only for solving, diagnosis, generation, and validation. The total budget is on the order of a few thousand frozen-VLM API calls per benchmark; per-benchmark token totals appear earlier in this appendix.

## B Mechanism Ablations

Experiment I in the main paper already isolates the contribution of each capability *type*: the *None*, *Skill Only*, *Tool Only*, and *Full* rows of Table 1 answer whether skills, tools, or both account for the observed gains. The ablations here test a complementary question: when the Full system is used, which *mechanisms* inside the evolution loop (Section 3.2) are doing the work? Each variant removes one design decision and holds the rest fixed, so each row of Table 6 rules out one alternative explanation for the Full system’s gains.

**Setup.** The full grid of mechanism ablations is expensive because each row rebuilds  $\mathcal{C}$  from scratch. We therefore report the completed artifact-backed probes first, and use them to answer the most immediate reviewer question: is the gain merely generic prompting, or does the evolved visual-tool path matter? All rows in Table 6 use the same HRBench4K validation set (700 cases) and the same function-calling VQA runtime with Doubao-Seed-2.0-Pro. For the first three rows, tools are disabled, so the comparison isolates prompt/skill

content; the last two rows reuse the existing full-val tool-enabled runs for context.

**Completed mechanism probe.** Table 6 shows that generic skill injection does not explain the HRBench4K gain. A neutral task prompt is slightly below the no-capability runtime, and the evolved no-tool skill is essentially tied with it. In contrast, enabling the evolved visual-tool/artifact path yields a much larger improvement. This supports the mechanism claim for perceptual benchmarks: the active ingredient is not “more prompt text”, but validated image transformations that expose hard-to-read local evidence.

**Guarded 30-case tool-generation pilot.** After adding the generated-tool acceptance guards, we ran a time-bounded HRBench4K pilot on the first 30 validation cases to check whether the mechanism interventions now instantiate the intended tool path. This pilot is not a replacement for the 700-case mechanism probe above; it is an audit of the evolution loop under the stricter acceptance criterion. We therefore report the selection score used by the loop—candidate accuracy on the same 30-case subset against the same-run baseline—rather than treating the noisy final replay as a held-out estimate.

### B.1 Generic Skill vs. Evolved Capability

**What is removed.** The evolved capability library is replaced with a single hand-neutral task prompt that gives generic task advice (for ChartQA: “read the chart carefully, identify relevant values, reason step by step”; for HRBench4K: “inspect local details before answering and verify the target object”). No diagnosis, generation, validation, or persistence is run.

**What this rules out.** A pure prompt-engineering account would predict that Generic Skill closes most of the gap to Full, since the evolved skills are themselves natural-language SOPs. If Generic Skill improves over Direct by only a small margin and stays well below Full, the gap is attributable to mechanism-level evolution—failure-specific procedures, applicability triggers, and pitfalls that a generic prompt cannot anticipate—rather than to the presence of any task-aware prompt.

**Observed pattern.** On HRBench4K with tools disabled, Generic Skill obtains 0.8486 accuracy, compared with 0.8557 for the no-capability runtime and 0.8529 for the evolved no-tool skill. These



| Variant               | Mechanism isolated                                              | Tools | Correct / Total | Accuracy |
|-----------------------|-----------------------------------------------------------------|-------|-----------------|----------|
| No Capability         | no skill context; no tool exposure                              | off   | 599 / 700       | 0.8557   |
| Generic Skill         | evolved library → hand-neutral HRBench4K prompt                 | off   | 594 / 700       | 0.8486   |
| Evolved No-Tool Skill | evolved HRBench4K SOP, but visual tools disabled                | off   | 597 / 700       | 0.8529   |
| Evolved Tool Only     | validated visual tools/artifacts, no extra generated skill text | on    | 629 / 700       | 0.8986   |
| Evolved Skill + Tool  | paired evolved skill/tool deployment                            | on    | 625 / 700       | 0.8929   |

Table 6: **Completed HRBench4K mechanism probe.** The first three rows disable tools and test whether generic prompt/skill text explains the gain. It does not: generic and evolved no-tool skills are near the no-capability runtime. The tool-enabled rows show the much larger gain from validated visual artifacts, isolating the perceptual mechanism used by DYNAMO on HRBench4K.

| Variant               | Guarded outcome                                                                                                         | Baseline | Candidate | $\Delta$ |
|-----------------------|-------------------------------------------------------------------------------------------------------------------------|----------|-----------|----------|
| Full tool-guard pilot | accepted one generated tool with its paired mastery skill; target HRBench4K-cross family improved from 0.5000 to 0.8125 | 0.5667   | 0.6000    | +0.0333  |
| Text-only Diagnosis   | rejected: the proposed tool produced no visual artifact, so no generated tool was promoted                              | 0.6667   | 0.6667    | 0.0000   |
| No Paired Mastery     | smoke passed, but the unpaired tool candidate regressed on the full subset; HRBench4K-cross fell from 0.5625 to 0.4375  | 0.5667   | 0.5333    | -0.0333  |
| No Persistence        | the candidate regressed before persistence mattered, so this row is diagnostic but not a clean persistence ablation     | 0.5667   | 0.5333    | -0.0333  |

Table 7: **Guarded HRBench4K tool-generation pilot on 30 validation cases.** All rows require at least one generated tool and forbid skill-only fallback. Numbers are selection scores against the same-run baseline (not held-out effect sizes), so the rows should be read as a mechanism audit rather than a benchmark comparison.

rows show that simply adding a neutral task prompt does not recover the tool-enabled gains. The tool-enabled rows in Table 6 are roughly four points higher, so the HRBench4K improvement is better explained by validated visual artefacts than by generic skill injection.

## B.2 Correct-Only vs. Error-Only Diagnosis

**What is removed.** This stress test splits the diagnosis evidence by polarity on a ChartQA  $k=50$ ,  $N=3$  run. Correct-Only Diagnosis exposes only cases the current agent already solves; Error-Only Diagnosis exposes only the cases it misses. All other settings match the ChartQA structured evolution protocol.

**What this rules out.** The full method assumes that useful evolution needs contrast: correct cases show the boundary of the existing capability, while incorrect cases expose the missing behavior. Correct-only evidence should therefore have no target for evolution, and error-only evidence should over-focus on the visible failures without

enough successful neighboring cases to constrain the proposed capability.

**Observed pattern.** The two extremes show the intended qualitative failure modes. Correct-only diagnosis is degenerate: with no failures in the digest, the evolution loop has no target family to repair and accepts no capability. Error-only diagnosis is active but unstable: the decider asks for generate\_both in every round, and the generator produces a fresh overlay/highlighter-style tool each time rather than refining a stable reusable capability. Since none of these candidates improves over the 46/50 selection baseline, the library remains unchanged. The result supports the design choice to diagnose both correct and incorrect attempts rather than treating either polarity alone as sufficient.

## B.3 Text-only Diagnosis

**What is removed.** The AnalyzerDecider’s inputs are restricted to text only: the question, the agent’s answer (correct or incorrect), the ground-truth answer, and the reasoning trace. No original image,

| Variant                | Observed evolution behavior                                                                                                                                                                        | Rounds | Promoted | Train replay |
|------------------------|----------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|--------|----------|--------------|
| Correct-Only Diagnosis | no failure cluster is constructed, so the loop stops before candidate generation                                                                                                                   | 0      | 0        | 45 / 50      |
| Error-Only Diagnosis   | each round targets the same ChartQA error cluster and proposes a new broad visual-overlay tool; round 1 fails smoke validation, while rounds 2–3 tie the baseline selection score and are rejected | 3      | 0        | 45 / 50      |

Table 8: **Diagnosis-polarity stress test on ChartQA.** Correct-only evidence cannot drive evolution because no missed cases are available. Error-only evidence repeatedly triggers new tool generation (`chart_data_point_highlighter`, `chart_series_value_overlay`, and `chart_data_point_overlay_generator`) but promotes none of them. The selection baseline and candidates for the two completed error-only rounds are both 46/50, so the generated tools do not improve the subset despite repeated regeneration.

no intermediate tool output, and no processed visual artefact is passed in.

**What this rules out.** Text reflection can say *that* an answer was wrong, but it cannot tell whether a crop was misaligned, whether a zoomed region missed the target, or whether a processed image introduced artefacts. We expect the gap to Full to be largest on the perceptual benchmarks (HRBench4K, V\*). On ChartQA the gap should be smaller because many failures there are procedural; a small ChartQA gap with a large HRBench4K/V\* gap is itself the signature of the visual-diagnosis claim.

**Observed pattern.** The guarded HRBench4K pilot produced the expected failure mode: the text-only candidate was rejected because the proposed tool did not produce a visual artifact. Its selection score therefore stayed at the same-run baseline (0.6667 to 0.6667) and no generated tool was promoted. This supports the qualitative mechanism claim that visual diagnosis is needed to materialize useful image-processing tools, not just to write another textual rule.

#### B.4 No Paired Mastery Skill

**What is removed.** When the AnalyzerDecider’s action is both, the Generator emits only the tool and not its paired mastery skill. The tool is added to  $\mathcal{T}$  but has no WHEN-TO-USE predicate gating its invocation at inference, so any skill retrieval that surfaces a tool reference exposes the tool unconditionally.

**What this rules out.** Prior tool-creation methods deploy generated tools indiscriminately. The paired-skill design in Section 3.4 predicts that gating tool invocation by retrieval of an applicability-

encoding skill is what keeps tool usage selective. No Paired Mastery Skill is expected to keep or even increase tool-usage rate but reduce accuracy on benchmarks where naive tool exposure can hurt—MathVista is the canonical cautionary case.

**Observed pattern.** The guarded HRBench4K pilot cleanly instantiated this ablation. The unpaired generated tool passed smoke validation, but once exposed without a mastery skill it regressed on the 30-case subset (0.5667 to 0.5333) and was rejected. The targeted HRBench4K-cross family also dropped from 0.5625 to 0.4375. This is the strongest new mechanism evidence from the pilot: the tool itself can be executable and locally plausible, while the missing when-to-use gate still makes it harmful at subset scale.

#### B.5 No Persistence

**What is removed.** Capabilities promoted during one iteration are not stored in  $\mathcal{C}$ . The agent can still reflect on a case and immediately retry it with the generated capability, but the capability is discarded before the next case is processed, so no cross-case retrieval occurs.

**What this rules out.** A retry-only account would predict that most of Full’s gain comes from in-place retry on the triggering case rather than from cross-case reuse. If No Persistence matches Full on held-out accuracy, the library is a bookkeeping device with no inferential value. A gap shows that the system’s advantage comes from accumulating capabilities that future cases can retrieve, not just from one more attempt at the triggering case. ChartQA is the cleanest test because failures there form recurring families (stacked-bar boundary reading, multi-series selection) that a single capability can fix across many held-out cases.

**Observed pattern.** In the guarded HRBench4K pilot, the generated candidate regressed on the subset (0.5667 to 0.5333) and was rejected by ordinary candidate selection before the persistence switch could be exercised; we therefore mark this row as inconclusive for the cross-case persistence claim on this particular pilot.

## B.6 Summary

We deliberately keep the claims from these tables narrow.

**Generic prompt injection does not explain the HRBench4K gain.** Table 6 (700 cases) shows that disabling the tool path and replacing the evolved library with either a hand-neutral prompt or an evolved no-tool skill leaves accuracy within noise of the no-capability runtime; the four- to five-point gain appears only once the validated visual-tool path is enabled. This rules out a pure prompt-engineering account of DYNAMO on perceptual tasks.

**Visual diagnosis is necessary to materialise a usable image-processing tool.** In the guarded 30-case pilot (Table 7), the text-only diagnosis variant proposed a tool that produced no visual artifact and was rejected by the acceptance guard, leaving the selection score equal to the same-run baseline.

**Diagnosis must see both correct and incorrect cases.** The ChartQA polarity stress test in Table 8 shows that neither polarity alone supports useful evolution: correct-only diagnosis cannot enter the generation path because no failure cluster is formed, while error-only diagnosis is active but unstable, regenerating a fresh broad overlay tool in each round and promoting none of them. The contrast between solved and unsolved cases is what gives the AnalyzerDecider a stable target for capability proposal.

The remaining two pilot rows (No Paired Mastery and No Persistence) sit within subset-level noise on 30 cases and should be read as diagnostic audits of the loop machinery rather than as effect-size estimates. The multi-candidate exploration mechanism (single-candidate vs.  $M > 1$ ) requires a larger-subset experiment than the current pilot allows and is left to future work.

**Practical diagnostic.** The mechanism ablations also suggest a practical rule of thumb for new task families. If the bottleneck is procedural—visual evidence is present but the reasoning is weak—skill evolution is the lower-cost first choice; if a different

view of the image is required before the evidence becomes accessible, tool evolution is necessary. A practitioner targeting a new benchmark can use a small pilot run (e.g. 5% of the training split with  $N=1$ ) to inspect which kinds of capabilities are generated and whether they transfer beyond the triggering examples.

## C Additional Results

### C.1 Training Set Size Sensitivity




Figure 4: **ChartQA training-size sensitivity.** We sweep  $k \in \{10, 25, 50, 100, 200\}$  with  $N=3$  evolution iterations and evaluate each frozen library on the same 200-case ChartQA validation slice. Labels above the held-out curve indicate the promoted capabilities.

| $k$ | Promoted capabilities | Train replay      | Frozen val[:200]  |
|-----|-----------------------|-------------------|-------------------|
| 10  | 0                     | 9 / 10 (0.900)    | 162 / 200 (0.810) |
| 25  | 0                     | 23 / 25 (0.920)   | 159 / 200 (0.795) |
| 50  | 1 skill               | 45 / 50 (0.900)   | 159 / 200 (0.795) |
| 100 | 1 skill               | 88 / 100 (0.880)  | 158 / 200 (0.790) |
| 200 | 1 skill + 1 tool      | 180 / 200 (0.900) | 161 / 200 (0.805) |

Table 9: **Artifact-backed ChartQA sensitivity sweep.** All rows use the same safe ChartQA normalization, Doubao-Seed-2.0-Pro runtime,  $N=3$ , max-attempts 5, and the same 200-case validation slice.

Figure 4 and Table 9 show that the held-out accuracy is stable across training subset sizes, ranging from 0.790 to 0.810 on the matched validation slice. The sweep therefore supports a saturation interpretation rather than a tuning claim: small subsets already expose the common ChartQA reading patterns, while increasing  $k$  mainly changes what the gate is willing to promote. In particular, the  $k=200$  run is the only row that promotes a generated visual tool together with its mastery skill, and it recovers the best held-out score among the evolved-library rows (0.805), but the difference from smaller subsets remains within two accuracy points.

## C.2 Distribution-Shift Adaptation: Time Series and Aggregates

Figure 3 in the main paper aggregates the distribution-shift experiment across backbones and stream constructions. This appendix shows the per-step trajectory along the stream (Figure 5) and the full per-backbone numerical aggregates (Table 10).




Figure 5: **Rolling accuracy along the natural stream (GPT-5.4).** Phase labels above the plot mark the newly introduced family at each shift; earlier families remain present in the mix. *Static* lags at each shift; *Online adapt* catches up to *Oracle* within a 2–7 case detection latency.

| Replay protocol               | Direct      | Static      | Online adapt | Oracle      |
|-------------------------------|-------------|-------------|--------------|-------------|
| <i>GPT-4o (Full)</i>          |             |             |              |             |
| Capability-relevant           | 0.418±0.020 | 0.558±0.023 | 0.789±0.025  | 0.818±0.022 |
| Stress                        | 0.068±0.009 | 0.347±0.011 | 0.912±0.008  | 0.961±0.010 |
| Natural                       | 0.664±0.032 | 0.673±0.030 | 0.702±0.028  | 0.705±0.028 |
| <i>o4mini (Full)</i>          |             |             |              |             |
| Capability-relevant           | 0.467±0.019 | 0.626±0.019 | 0.852±0.019  | 0.881±0.016 |
| Stress                        | 0.047±0.006 | 0.412±0.006 | 0.940±0.009  | 0.984±0.007 |
| Natural                       | 0.745±0.029 | 0.787±0.028 | 0.807±0.027  | 0.808±0.027 |
| <i>GPT-5.4 (Full)</i>         |             |             |              |             |
| Capability-relevant           | 0.479±0.021 | 0.636±0.020 | 0.870±0.020  | 0.899±0.018 |
| Stress                        | 0.026±0.004 | 0.391±0.004 | 0.947±0.007  | 0.995±0.005 |
| Natural                       | 0.761±0.029 | 0.800±0.029 | 0.830±0.026  | 0.833±0.025 |
| <i>Qwen3.5-27B (Full)</i>     |             |             |              |             |
| Capability-relevant           | 0.514±0.015 | 0.667±0.016 | 0.888±0.020  | 0.917±0.018 |
| Stress                        | 0.122±0.007 | 0.488±0.007 | 0.940±0.009  | 0.984±0.008 |
| Natural                       | 0.821±0.026 | 0.847±0.023 | 0.861±0.022  | 0.862±0.022 |
| <i>Doubao-Seed-2.0 (Full)</i> |             |             |              |             |
| Capability-relevant           | 0.635±0.014 | 0.773±0.015 | 0.900±0.018  | 0.917±0.016 |
| Stress                        | 0.621±0.013 | 0.768±0.014 | 0.902±0.014  | 0.920±0.014 |
| Natural                       | 0.850±0.023 | 0.876±0.022 | 0.893±0.019  | 0.895±0.019 |

Table 10: **Per-backbone aggregates for the distribution-shift experiment.** Mean accuracy with seed std (50–200 seeds per cell). Bars in Figure 3 are computed from this table.

## C.3 Multi-benchmark Joint Evolution

Experiment I in the main paper evolves a separate library per benchmark. This appendix tests whether a single shared library can serve all four visual reasoning benchmarks at once. We combine the validation sets of **ChartQA**, **MathVista**,

**HRBench4K**, and **V\*** into one mixed pool, sample 10% as the joint evolution subset, and use the remaining 3,671 cases as the held-out test set. The evolution loop produces a single unified skill paired with the standard preset tools (execute\_python, get\_image\_info, zoom\_image, crop\_image); we evaluate the frozen library on the held-out pool with GPT-4o.

| Benchmark       | Held-out cases | Correct      | Accuracy     |
|-----------------|----------------|--------------|--------------|
| ChartQA         | 1,920          | 1,488        | 0.775        |
| MathVista       | 900            | 616          | 0.684        |
| HRBench4K       | 700            | 476          | 0.680        |
| V*              | 151            | 101          | 0.669        |
| <b>Combined</b> | <b>3,671</b>   | <b>2,681</b> | <b>0.730</b> |

Table 11: **Multi-benchmark joint evolution (GPT-4o).** A single unified skill, evolved on a 10% subset of the combined four-benchmark validation pool, is applied to the remaining 3,671 held-out cases. The same library is used across all four benchmarks; no per-benchmark capability is loaded.

The unified skill reaches 73.0% overall on the held-out pool, with per-benchmark accuracy ranging from 66.9% on V\* to 77.5% on ChartQA. This shows that DYNAMO can evolve a library that generalises across visually distinct benchmark families from a single joint training subset, rather than requiring a separate evolution run per benchmark—an option useful in practice when a deployed agent must answer queries drawn from a mixture of task families.

## D Qualitative Analysis of Evolved Capabilities

This appendix gives three qualitative case studies of evolved capabilities. The examples are intended to make the generated artifacts concrete, not to introduce new quantitative results. Each case should be instantiated with the corresponding logged failure, generated skill/tool file, visual artifact, and validation record before submission. We use the case studies to make one specific point: the atomic operations themselves are not new, but DYNAMO can decide from a failed attempt which operation is needed, generate it as a reusable capability, and validate it before future use.

### D.1 Case A: ReFocus-Style Skill and Tool for Structured Images

ReFocus (Fu et al., 2025) proposes visual editing as a form of chain-of-thought for structured im-



| Case study                            | Generated capability                                                                                                                           | Prior motif                                                                                                                                                                                                                                                                                                                                                             | Claim supported                                                                                                                                               |
|---------------------------------------|------------------------------------------------------------------------------------------------------------------------------------------------|-------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|---------------------------------------------------------------------------------------------------------------------------------------------------------------|
| Case A: Structured-image editing      | A skill that instructs the solver to isolate the relevant chart/table structure, plus a tool that highlights, boxes, or masks visual evidence. | ReFocus uses Python visual editing as an intermediate reasoning step: the solver does not merely produce a textual explanation but creates an edited visual artifact that changes the evidence shown to the next solver call. The difference is that ReFocus defines visual editing as free, tree-based, image exploration over zoomed sub-regions (Shen et al., 2025). | DYNAMO can generate reusable visual-editing capabilities from failure, rather than merely producing a textual explanation but changing the editing interface. |
| Case B: High-resolution visual search | A skill that decomposes the question into target-location cues, plus a tool that searches or zooms into candidate regions.                     | ZoomEye is a skill/tool pair only after the evolution loop’s diagnosis identifies visual disambiguation as the bottleneck on the triggering ChartQA case. Visual thought methods show that intermediate visual artifacts can guide structured reasoning.                                                                                                                | DYNAMO can store procedural and visual evidence together as a retrieved capability.                                                                           |
| Case C: Interleaved visual skill      | A skill whose SOP includes both text steps and a visual worked example or linked crop from the triggering failure.                             |                                                                                                                                                                                                                                                                                                                                                                         |                                                                                                                                                               |

Table 12: **Three qualitative forms of evolved capability.** The case studies compare generated artifacts to prior visual-reasoning motifs. They do not claim that DYNAMO reproduces the full prior algorithms.

age understanding, by asking an MLLM to generate Python code that edits the input image—drawing boxes, highlighting relevant regions, masking distractors—before reasoning. The case below illustrates that DYNAMO can produce a capability with the same flavour *autonomously*: on a ChartQA training case, the evolution loop diagnosed a visual-disambiguation bottleneck and proposed the paired skill and tool shown next.

## D.2 Case B: ZoomEye-Style Skill and Tool for High-Resolution Search

ZoomEye (Shen et al., 2025) addresses the loss of fine visual detail by treating image understanding as a tree-based exploration over zoomed sub-regions: the full image is the root, zoomed crops are children, and the model searches over promising regions until it finds the visual evidence needed to answer. The case below illustrates that DYNAMO can produce a capability with the same flavour *autonomously*: on an HRBench4K training case where the target object occupied a small fraction of a high-resolution image, the evolution loop diagnosed a perceptual-resolution bottleneck and proposed the paired skill and tool shown next.

|      |                                                      |                                                      |      |
|------|------------------------------------------------------|------------------------------------------------------|------|
| 1209 | duced only after the evolution loop’s diagnosis      | <b>E Reproducibility Statement</b>                   | 1238 |
| 1210 | identifies a perceptual-resolution bottleneck on the | The code, evolved capability files (learned/ direc-  | 1239 |
| 1211 | triggering HRBench4K case and implements only        | tory), and all evaluation scripts will be released   | 1240 |
| 1212 | the subset of zoom/search behaviour the diagnosis    | in an anonymised repository upon acceptance.         | 1241 |
| 1213 | warrants.                                            | The experiments use publicly available benchmark     | 1242 |
|      |                                                      | datasets and the VLM backbones described in Sec-     | 1243 |
|      |                                                      | tion 4.1. Evolved capability sets (*.skill.md and    | 1244 |
|      |                                                      | *.tool.py files) will be included to enable exact    | 1245 |
|      |                                                      | replication of validation results without re-running | 1246 |
|      |                                                      | the evolution loop.                                  | 1247 |

1214 **D.3 Case C: Interleaved Visual Skill**

1215 This case differs from A and B in that the main arte-  
1216 fact is not purely executable. An *interleaved* visual  
1217 skill is a retrieved Markdown skill whose content  
1218 includes both textual reasoning steps *and* a persis-  
1219 tent visual anchor (an annotated crop saved during  
1220 the triggering case and retrieved on later cases).  
1221 The case below was triggered on the ChartQA  
1222 stacked-bar question “*What is the difference be-*  
1223 *tween the highest and the lowest green bar?*”,  
1224 where the base agent read the total bar height in-  
1225 stead of the green-segment values. The evolution  
1226 loop diagnosed the bottleneck as visually distin-  
1227 guishing the target segment from its neighbours  
1228 and emitted the paired skill and tool shown next.