

DYNAMO: Dynamic Skill-Tool Evolution for Vision-Language Agents

Anonymous ACL submission

Abstract

Improving vision-language models (VLMs) on visual reasoning typically requires retraining or hand-designed prompts and tools. We present DYNAMO, a training-free framework in which a frozen VLM agent inspects its own correct and incorrect attempts on a small training subset and evolves two complementary capabilities: reusable reasoning skills for cognitive bottlenecks, and executable visual tools—gated by learned applicability—for perceptual ones, accumulating them in a persistent library without weight updates. Across four visual reasoning benchmarks and six VLM backbones, DYNAMO improves direct inference in 21 of 23 settings (avg. +4.8 acc). The same framework extends to the tool-mastery setting on GTA, where mastery profiling lifts step-level tool selection across all tested backbones. Against task-specific RL (VTool-R1, DeepEyes), DYNAMO matches or approaches RL accuracy at a fraction of the compute (60–94% gap recovery on chart/table; 91.1 vs. 90.1 on V*-7B) and combines additively with RL when available.

1 Introduction

Vision-language models (VLMs) have improved rapidly, yet adapting a frozen VLM to a new visual-reasoning task family still typically requires manually curated SFT data or a hand-designed RL pipeline—per task and per backbone. We ask whether the agent can build its own task-family-specific capability set instead, by inspecting its own behaviour on a small training subset and without any weight updates.

We instantiate this idea in DYNAMO, a training-free framework that evolves two complementary capability types and dynamically equips them at inference. **Skills** are structured Markdown SOPs for *cognitive* bottlenecks, where the visual evidence is available but the agent’s reasoning procedure is weak. **Tools** are short Python programs for *perceptual* bottlenecks, where the agent needs a trans-




Figure 1: **Labor-free task adaptation via skill-tool evolution.** Per-task VLM adaptation typically requires manually curated SFT data with a hand-designed RL pipeline for every new task family. DYNAMO instead lets a frozen VLM agent evolve its own skill and tool set per task family without human supervision, and dynamically equips the matching set when a case from that family arrives.

formed view—a crop, zoom, contrast adjustment, or rerendering—before the evidence becomes legible. On each iteration over a sampled sub-training set, DYNAMO diagnoses both the correct and the incorrect attempts using their images and reasoning traces, proposes multiple candidate skill+tool combinations, and promotes the highest-validation-accuracy candidate into a persistent library. A mastery phase further learns when each tool is safe to invoke, so that promoted capabilities are deployed selectively rather than indiscriminately.

We evaluate three questions. (I) Does DY-

NAMO work across diverse benchmarks and backbones? On ChartQA, MathVista, HRBench, and V* across six VLM backbones, skill-tool co-evolution improves direct inference in 21 of 23 model–benchmark settings, averaging +4.8 accuracy points. **(II) Does the framework extend to a tool-mastery setting with a curated tool set?** On GTA, mastery profiling lifts step-level tool selection, argument, and instruction-following accuracy across all tested backbones. **(III) Can DYNAMO match task-specific RL across different task families?** Across three task families—chart, table, and high-resolution perception—we compare DYNAMO to a representative RL method on each (VTool-R1 for chart/table; DeepEyes for perception). DYNAMO matches or approaches RL accuracy at a fraction of the compute (60–94% gap recovery on chart/table; 91.1 vs. 90.1 on V*-7B perception), and combines additively with RL when both are available.

Contributions.

- (1) A **problem formulation** that recasts per-task VLM adaptation as labor-free capability evolution: a frozen agent builds its own task-family-specific skill and tool set from a small training subset, replacing manual SFT data curation and task-specific RL pipelines.
- (2) A **multi-candidate evolution loop with mastery** that diagnoses both correct and incorrect attempts, generates candidate skill+tool combinations, selects the best by validation accuracy, and learns per-tool applicability boundaries.
- (3) **Empirical evidence** across four visual-reasoning benchmarks, six VLM backbones, the GTA tool-mastery benchmark, and three task families (chart, table, perception) each compared against its representative RL method, showing that capability evolution is a viable training-free alternative—and add-on—to task-specific RL.

2 Related Work

Self-improving agents. A growing line of work lets LLM agents improve from experience without gradient updates: Reflexion (Shinn et al., 2023) keeps a verbal reflection buffer that resets between episodes; Voyager (Wang et al., 2024a) and ExpeL (Zhao et al., 2024) accumulate persistent skill libraries in game and tool-use domains; AutoManual (Chen et al., 2024) and EvolveR (Anonymous,

2025a) refine instruction manuals or reasoning strategies via self-play. All operate in text, code, or game domains, never generating executable visual processing nor deciding whether a bottleneck is cognitive or perceptual. DYNAMO extends the persistent-library paradigm into the visual domain with a multimodal diagnostic that text-only frameworks cannot perform.

Tool creation for language models. A complementary line generates new tools on demand: CREATOR (Qian et al., 2023) prompts an LLM to write Python helpers for problems it cannot solve directly; LATM (Cai et al., 2023) and ToolLLM (Qin et al., 2024) scale this to large API collections. These methods generate tools in textual or mathematical domains and do not profile when a generated tool succeeds or fails—a gap that matters for visual tools, where the same operation can help on most images but hurt on specific patterns (e.g. a sharpening filter degrading low-contrast charts). DYNAMO’s mastery phase profiles each tool on a held-out case set to characterise its applicability before promotion.

Visual reasoning with tools. Several recent systems equip VLMs with visual processing tools. VTool-R1 (Anonymous, 2025d), DeepEyes (Zheng et al., 2025), Adaptive-CoF (Zhang et al., 2025), PixelReasoner (Anonymous, 2025b), and V-Thinker (Anonymous, 2025c) all rely on SFT and/or RL with curated trajectories to teach a VLM when and how to invoke a fixed tool inventory. ZoomEye (Shen et al., 2024) and ReFocus (Fu et al., 2025) are training-free but assume a fixed zoom/search algorithm or visual-editing interface. DYNAMO instead generates and validates the skills and tools themselves from the agent’s own behaviour on a small training subset, with no weight updates and no predefined tool inventory.

Tool-augmented VLMs with predefined tool sets. VisProg (Gupta and Kembhavi, 2023), ViperGPT (Surís et al., 2023), and Chameleon (Lu et al., 2023) compose hand-curated tool libraries (detectors, OCR, arithmetic modules) via program synthesis. They handle broad compositional visual tasks effectively but cannot recover when the required tool is missing from the human-designed inventory. DYNAMO fills this gap by generating new tools when the current set is insufficient, bridging fixed-tool composition and training-based tool learning.




Figure 2: **DYNAMO evolution loop.** On a sampled sub-training set, DYNAMO diagnoses both correct and incorrect attempts, explores candidate skill+tool capabilities, and promotes the best by validation into a persistent library.

3 Method

3.1 Problem Formulation

Let $\mathcal{D}_{\text{train}} = \{(x_i, y_i, \mathbf{v}_i)\}_{i=1}^k$ denote a small training subset of k cases, where x_i is the question, y_i the ground-truth answer, and $\mathbf{v}_i$ the associated visual input (one or more images). Let π_θ be a frozen VLM backbone. We define a capability set $\mathcal{C} = (\mathcal{S}, \mathcal{T})$, where $\mathcal{S}$ is a library of *skills* (structured reasoning SOPs) and $\mathcal{T}$ is a library of *tools* (executable Python programs). The agent decomposes solving into a retrieval step and a reasoning step:

$$f_{\mathcal{C}}(x, \mathbf{v}; \pi_\theta) = \pi_\theta(\cdot \mid x, \mathbf{v}, \text{Retrieve}(x, \mathbf{v}; \mathcal{C})), \quad (1)$$

where Retrieve returns the subset of skills and tools relevant to the current input. Given a held-out validation set $\mathcal{D}_{\text{val}}$, our goal is to learn $\mathcal{C}$ that maximises agent accuracy without any weight update:

$$\mathcal{C}^* = \arg \max_{\mathcal{C}} \text{Acc}(f_{\mathcal{C}}; \mathcal{D}_{\text{val}}), \quad (2)$$

where the maximisation is over $\mathcal{C}$ only and π_θ remains frozen ($\nabla_\theta = 0$).

3.2 Evolution Loop

Each evolution iteration over $\mathcal{D}_{\text{train}}$ has three phases (Figure 2).

Diagnose. The agent (Eq. 1) solves $\mathcal{D}_{\text{train}}$ with the current $\mathcal{C}$, then samples a sub-training set

$\mathcal{D}_{\text{sub}} \subseteq \mathcal{D}_{\text{train}}$ containing both correct and incorrect attempts. The ANALYZERDECIDER inspects $\mathcal{D}_{\text{sub}}$ —the questions, original images $\mathbf{v}_i$, reasoning traces τ_i , and any intermediate tool outputs—and emits a root-cause analysis grounded in visual evidence, an identified bottleneck type (*cognitive* when the visual evidence is available but the reasoning procedure is weak, or *perceptual* when the agent needs a transformed visual input), and an action $a \in \{\text{skill}, \text{tool}, \text{both}\}$. Visual input here is essential: deciding whether a crop is misaligned, whether labels are legible, or whether a processed image introduced artefacts requires inspecting the image, not just the reasoning trace.

Explore. Conditioned on the diagnosis, the GENERATOR proposes M candidate skill+tool combinations, each addressing the diagnosed bottleneck with a different strategy or implementation.

Validate and promote. Each candidate is evaluated on $\mathcal{D}_{\text{train}}$ and the highest-accuracy combination is promoted:

$$m^* = \arg \max_{m \in \{1, \dots, M\}} \text{Acc}(f_{\mathcal{C} \cup (s_m, t_m)}; \mathcal{D}_{\text{train}}), \quad (3)$$

$$\mathcal{C} \leftarrow \mathcal{C} \cup \{(s_{m^*}, t_{m^*})\}.$$

Eq. 3 approximates the validation objective in Eq. 2 by its training-set counterpart, and subsumes both an *origin* requirement (the new capability must improve over the previous $f_{\mathcal{C}}$ on the cases it targets)

and a *regression* requirement (it must not degrade currently-correct cases). When the action is both, the paired skill additionally serves as the tool’s mastery SOP, gating tool invocation at inference (Section 3.4). The full procedure runs for N iterations (default $N=3$); pseudocode is given in Appendix A (Algorithm 1).

3.3 Skills

A *skill* is a structured Markdown document with four fields: **When-to-Use** (a trigger predicate over question and image type), **Strategy** (a numbered step-by-step SOP), **Common Pitfalls**, and **Worked Example**. The Generator writes a skill conditioned on the AnalyzerDecider’s root-cause analysis; if a skill covering the same problem class already exists in $\mathcal{S}$, new insights are merged into it rather than creating a duplicate, preventing library bloat. At inference, the Solver retrieves the top- K skills by BM25 similarity to the question and appends them to the system prompt.

3.4 Tools and Mastery

A *tool* is a Python function (≤ 150 lines) that takes one or more image paths and returns processed images or extracted data, using standard CV libraries (OpenCV, PIL, NumPy). The Generator writes a tool conditioned on a perceptual bottleneck diagnosis; representative examples produced in our experiments include chart re-rendering with enlarged axis labels, saliency-guided sub-region extraction, and contrast enhancement for low-quality documents.

Mastery as a paired skill. When the AnalyzerDecider’s action is both, the Generator emits a skill alongside the tool, and this paired skill is the tool’s mastery SOP. Its When-to-Use predicate specifies the input patterns on which the tool helps, its Common Pitfalls flag the patterns on which it hurts, and its Strategy instructs the Solver how to invoke the tool and consume its output. Because tool invocation is gated by retrieving the paired skill, deployment is selective by construction rather than indiscriminate—a property that distinguishes DYNAMO from prior tool-creation methods (Qian et al., 2023; Cai et al., 2023) that expose generated tools without learned applicability boundaries.

External tool sets (Mode B). When a curated tool set $\mathcal{T}_0 \neq \emptyset$ is already provided, DYNAMO can operate in **Mode B**: skip tool generation and instead synthesise a mastery skill for each provided tool $t_j \in \mathcal{T}_0$, learning when to invoke it from the

agent’s own behaviour on $\mathcal{D}_{\text{train}}$. Mode B requires no code generation, making it practical when tool quality is already high but deployment strategy is unknown. Experiment II (Section 4.3) evaluates this mode.

3.5 Online Adaptation to Distribution Shift

In real deployments, DYNAMO runs against a streaming query distribution that can shift over time—a stream may begin with high-resolution perception cases and later start including MathVista-style reasoning cases, or vice versa. We design DYNAMO to detect such shifts and *update* its skills and tools online, without weight updates and without the practitioner having to anticipate the shift or prepare a per-family training set.

DYNAMO continuously samples a sub-training set from the most recent queries and runs the evolution loop of Section 3.2 on them, accumulating skills and tools into the library $\mathcal{C}$. Let W be a rolling-window length. The dominant task family at step t is

$$\hat{f}_t = \arg \max_f \sum_{i=t-W+1}^t \mathbf{1}[\text{fam}(x_i) = f], \quad (4)$$

where $\text{fam}(\cdot)$ classifies each query from its question and image features. When Eq. 4 flips ($\hat{f}_t \neq \hat{f}_{t-1}$), the next evolution iteration is triggered on the new sub-stream and produces updated skills and tools that handle the new distribution. The library is thus kept in sync with the current query distribution without any human intervention. We treat this autonomous online adaptation as one of the method’s contributions; Experiment I (Section 4.2) evaluates it against a static baseline that never refreshes and an oracle whose full capability set is available from the start.

4 Experiments

4.1 Experimental Setup

Benchmarks. For the evolution-from-scratch setting (Exp I) we use four visual reasoning benchmarks: **ChartQA** (Masry et al., 2022) (multi-step numerical reasoning over charts; 2,500 test questions, relaxed accuracy with 5% tolerance), **MathVista** (Lu et al., 2024) (math reasoning grounded in figures and plots; 1,000 testmini questions, accuracy), **HRBench** (Anonymous, 2024) (perception on $\geq 4K$ -resolution images; accuracy), and **V*** (Wu and Xie, 2023) (visual search for an object’s property; ≈ 500 questions, accuracy); the first two stress

cognitive bottlenecks while the latter two stress *perceptual* ones. For the tool-mastery setting (Exp II) we use **GTA** (Wang et al., 2024b), a benchmark of 229 real-world tasks with tools across perception, operation, logic, and creativity categories. GTA reports step-by-step accuracy on tool selection, argument prediction, and instruction following, plus an end-to-end answer accuracy. For the RL comparison (Exp III) we follow the VTool-R1 protocol (Anonymous, 2025d) on its **Chart** and **Table** reasoning splits and reuse **V*** and **HRBench4K** as the high-resolution splits of the DeepEyes protocol (Zheng et al., 2025).

Foundation models. Experiment I evaluates five proprietary and open-weight VLMs: **GPT-4o**, **o4-mini**, **GPT-5.4**, **Doubao-Seed-2.0**, and **Qwen3.5-27B**. Experiment II evaluates **GPT-4o**, **GPT-5.4**, and **Doubao-Seed-2.0** on the GTA benchmark, plus a controlled **Doubao-Seed-2.0-Pro** analysis that inspects how provided-tool skills change with tool abstraction and task environment. Experiment III uses the **Qwen2.5-VL** family (3B, 7B, 32B) to align with the VTool-R1 evaluation protocol. All VLM weights are *frozen* throughout; no gradient updates are performed.

Evolution protocol. For each benchmark, we sample 20% of the official training split as $\mathcal{D}_{\text{train}}$ and run the DYNAMO evolution loop for $N=3$ iterations. The capability set $\mathcal{C}$ is then frozen and evaluated on the held-out validation / test split. We report the average over three independent random seeds (different 20% subsets) to account for sampling variance where repeated runs are available.

Baselines. For Experiment I, we compare four evolution configurations that isolate the contribution of each component: **None (Base Agent)** invokes the frozen VLM with no evolved capabilities (the zero-shot baseline); **Skill Only** restricts the Generator to producing reasoning skills (no tool generation); **Tool Only** restricts the Generator to producing executable visual tools (no skill generation); **Full** is the unrestricted DYNAMO system, where the AnalyzerDecider chooses on each case whether to generate a skill, a tool, or both.

Evaluation metrics. All primary metrics are task-level accuracy on the held-out split. For Experiment II we report four of GTA’s five canonical metrics: three step-by-step metrics—**InstAcc** (Instruction-Following Accuracy: fraction of steps

executed without errors), **ToolAcc** (Tool Selection Accuracy), and **ArgAcc** (Argument Prediction Accuracy)—and the end-to-end **AnsAcc** (Answer Accuracy on full tool-using execution). For the supplementary controlled GTA analysis, we also report tool usage rate, average number of tool calls, and dominant tool-call chains. Experiment III follows the VTool-R1 evaluation protocol (Anonymous, 2025d) on Chart and Table reasoning splits, and uses DeepEyes (Zheng et al., 2025) as the RL reference for high-resolution **V*** and **HRBench4K** visual-search tasks.

4.2 Experiment I: Autonomous Evolution from Scratch

We test whether DYNAMO can bootstrap a useful capability library from scratch ($\mathcal{S}=\emptyset$, $\mathcal{T}=\emptyset$) across five backbones and four benchmarks, and whether the library remains useful when the query distribution shifts over time (Section 3.5).

Per-benchmark results. Table 1 reports accuracy for the four configurations on each backbone. Comparing the Full system to the zero-shot *None* baseline, DYNAMO improves direct inference on all 20 model–benchmark cells, with an average gain of +4.7 accuracy points. The largest gains are on o4-mini / HRBench (+12.7), GPT-5.4 / HRBench (+10.7), Doubao-Seed-2.0 / V* (+8.8), and Qwen3.5-27B / HRBench (+8.2). The relative strength of *Skill Only* and *Tool Only* also follows the cognitive / perceptual split: Skill Only is competitive on ChartQA and MathVista, where the gain comes from better decomposition and calculation procedures, while Tool Only and Full dominate on HRBench and V*, where sub-region inspection and visual search are central. On a few cells (GPT-4o / MathVista, Doubao / V*) Skill Only or Tool Only slightly edges out Full—an expected consequence of leaving capability choice to the AnalyzerDecider rather than enforcing both per case.

Capability quality. Appendix D inspects three representative forms of evolved capability: a ReFocus-style chart/table editing case, a ZoomEye-style coarse-to-fine search case, and an interleaved skill that pairs a textual SOP with visual evidence. The case studies illustrate what DYNAMO generates; they do not claim to reproduce the full algorithms of ReFocus (Fu et al., 2025) or ZoomEye (Shen et al., 2024).

Model	Evolution Strategy	HRBench	MathVista	V*	ChartQA
GPT-4o	None (Base Agent)	0.6386	0.6711	0.5828	0.7552
	Skill Only	0.6733±0.0121	0.7122±0.0038	0.6424±0.0066	0.7785±0.0026
	Tool Only	0.6452±0.0044	0.6900±0.0029	0.6225±0.0066	0.7722±0.0029
	Full	0.6886±0.0071	0.7189±0.0029	0.6689±0.0239	0.7799±0.0045
o4-mini	None (Base Agent)	0.73	0.753333	0.72847	0.7833
	Skill Only	0.7724±0.0058	0.7470±0.0107	0.7616±0.0066	0.7866±0.0151
	Tool Only	0.7457±0.1670	0.7356±0.0328	0.8387±0.0629	0.7920±0.0213
	Full	0.8571±0.0099	0.7559±0.0055	0.8698±0.0076	0.8007±0.0081
gpt5.4	None (Base Agent)	0.7471	0.7589	0.7748	0.7974
	Skill Only	0.7805±0.0036	0.7752±0.0039	0.7682±0.0199	0.8090±0.0022
	Tool Only	0.8367±0.0346	0.7663±0.0061	0.8477±0.0132	0.8108±0.0052
	Full	0.8538±0.0036	0.7815±0.0101	0.8521±0.0233	0.8276±0.0278
Doubao-Seed-2.0	None (Base Agent)	0.8214	0.8544	0.8675	0.8152
	Skill Only	0.8600±0.0049	0.8556±0.0109	0.8609±0.0115	0.8198±0.0031
	Tool Only	0.9005±0.0022	0.8659±0.0061	0.9448±0.0038	0.8156±0.0068
	Full	0.8990±0.0064	0.8681±0.0105	0.9558±0.0038	0.8347±0.0295
Qwen3.5-27B	None (Base Agent)	0.8171	0.8088	0.88079	0.8114
	Skill Only	0.8721±0.0030	0.8152±0.0080	0.8675±0.0066	0.8097±0.0090
	Tool Only	0.8986±0.0000	0.8241±0.0034	0.9514±0.0101	0.8115±0.0041
	Full	0.9010±0.0008	0.8252±0.0023	0.9603±0.0175	0.8158±0.0026

Table 1: **Autonomous evolution from scratch.** Accuracy across five VLM backbones and four benchmarks. The Full DYNAMO configuration improves over the zero-shot *None* baseline on all 20 model–benchmark cells.

Online adaptation to distribution shift. To simulate real online user traffic—where the task distribution shifts over time as different users and requests arrive—we replay a non-stationary stream that begins on one family (HRBench) and then drifts into mixtures dominated by MathVista, ChartQA, and finally V*. The replay reuses completed per-case outputs, isolating the adaptation policy from the cost of regenerating capabilities online. We compare *Direct* (no capabilities), *Static* (capabilities evolved on the initial distribution only and never updated), *Online adapt* (DYNAMO detects the shift and updates skills and tools on the new sub-stream), and *Oracle* (capabilities pre-evolved on every family, an upper bound), each on three stream constructions: *Capability-relevant* (cases whose outcome depends on the matching capability), *Stress* (cases the matching capability fixes), and *Natural* (fixed-seed mixture).

Table 2 shows *Online adapt* consistently lands closer to *Oracle* than to *Static*. On GPT-5.4 capability-relevant it reaches 0.870 versus 0.636 static and 0.899 oracle; on stress, *Direct* collapses to 0.026, *Static* reaches 0.391, and *Online adapt* recovers to 0.947. On the natural mix—where many cases do not require a newly-evolved capability—the gain is smaller but still positive (0.830 vs. 0.800

Replay protocol	Direct	Static	Online adapt	Oracle
<i>GPT-4o (Full)</i>				
Capability-relevant	0.418±0.020	0.558±0.023	0.789±0.025	0.818±0.022
Stress	0.068±0.009	0.347±0.011	0.912±0.008	0.961±0.010
Natural	0.664±0.032	0.673±0.030	0.702±0.028	0.705±0.028
<i>o4mini (Full)</i>				
Capability-relevant	0.467±0.019	0.626±0.019	0.852±0.019	0.881±0.016
Stress	0.047±0.006	0.412±0.006	0.940±0.009	0.984±0.007
Natural	0.745±0.029	0.787±0.028	0.807±0.027	0.808±0.027
<i>GPT-5.4 (Full)</i>				
Capability-relevant	0.479±0.021	0.636±0.020	0.870±0.020	0.899±0.018
Stress	0.026±0.004	0.391±0.004	0.947±0.007	0.995±0.005
Natural	0.761±0.029	0.800±0.029	0.830±0.026	0.833±0.025
<i>Qwen3.5-27B (Full)</i>				
Capability-relevant	0.514±0.015	0.667±0.016	0.888±0.020	0.917±0.018
Stress	0.122±0.007	0.488±0.007	0.940±0.009	0.984±0.008
Natural	0.821±0.026	0.847±0.023	0.861±0.022	0.862±0.022
<i>Doubao-Seed-2.0 (Full)</i>				
Capability-relevant	0.635±0.014	0.773±0.015	0.900±0.018	0.917±0.016
Stress	0.621±0.013	0.768±0.014	0.902±0.014	0.920±0.014
Natural	0.850±0.023	0.876±0.022	0.893±0.019	0.895±0.019

Table 2: **Online adaptation to distribution shift.** Replay accuracy on a non-stationary user-traffic stream. *Online adapt* stays close to the *Oracle* upper bound and far above *Static* across all five backbones.

static). The other four backbones follow the same pattern.

4.3 Experiment II: Learning to Use Pre-Provided Tools

Results on GTA. Table 3 compares *Base Agent* (zero-shot VLM with the provided GTA tools ex-

Model	Method	ToolAcc	ArgAcc	InstAcc	AnsAcc
GPT-4o	Base Agent	80.3	62.0	52.0	66.1
	Skill Evolution (Ours)	85.3	67.4	57.3	57.9
GPT-5.4	Base Agent	53.0	36.9	33.3	66.1
	Skill Evolution (Ours)	85.7	63.4	59.9	68.6
Doubao-Seed-2.0	Base Agent	67.7	45.5	41.6	76.9
	Skill Evolution (Ours)	86.4	61.3	57.7	81.8

Table 3: **Learning to use pre-provided tools on GTA.** DYNAMO Mode B (skills + mastery only; no new tool generation) improves all step-level metrics on every backbone, and improves end-to-end answer accuracy on the two stronger backbones.

posed) to DYNAMO Mode B (skills + mastery learned on the training set, no new tool generation) across three backbones. Mode B improves all three step-level metrics on every backbone: **InstAcc**, **ToolAcc**, and **ArgAcc** each gain double-digit points on GPT-5.4 (+27, +33, +27) and Doubao-Seed-2.0 (+16, +19, +16), and gain +5 on GPT-4o. End-to-end **AnsAcc** also improves on the two stronger backbones (GPT-5.4 +2.5, Doubao +4.9). The only regression is GPT-4o **AnsAcc** (−8.2): the evolved mastery SOP recommends tool invocation in a broader set of conditions than GPT-4o can reliably execute, leading it to invoke tools on cases it would have answered correctly without them. This is consistent with the Exp I observation that evolution gains depend on the backbone’s instruction-following capacity.

Controlled tool-policy analysis. Table 4 runs a complementary controlled study on GTA with Doubao-Seed-2.0-Pro to inspect what the learned policy actually does when either the tool abstraction or the task environment changes. Two findings stand out. First, fixing the environment to OCR+Calculator, stronger composite tools compress the policy: atomic tools reach 86.7 via an OCR→Calculator chain (22/30 cases), while VisualArithmeticSolver reaches 90.0 with a one-call policy (30/30). Second, the same OCR tool plays different roles in different environments: in OCR+Calc it feeds arithmetic (OCR→Calculator, 19/20), while in OCR+Search it feeds external lookup (OCR→Search→OCR→Search, 20/20).

4.4 Experiment III: RL and Self-Evolution Across Visual Tool Tasks

We compare DYNAMO against task-specific RL on three task families using the Qwen2.5-VL family (3B, 7B, 32B): Chart and Table reasoning under the VTool-R1 protocol (Anonymous, 2025d), and high-

Study	Condition	Cases	Acc.	Tool use	Avg. calls
Tool strength	Direct, no tools	30	83.3	0.0	0.00
Tool strength	Atomic OCR+Calc	30	86.7	100.0	1.83
Tool strength	Composite VAS	30	90.0	100.0	1.00
OCR role	OCR+Calc env	20	80.0	100.0	2.20
OCR role	OCR+Search env	20	85.0	100.0	2.35

Table 4: **Controlled tool-policy analysis on GTA (Doubao-Seed-2.0-Pro).** Top: stronger composite tools compress the policy at higher accuracy. Bottom: the same OCR tool exhibits different downstream behavior in different environments (dominant call chains discussed in prose). Acc./Tool use in %; Avg. calls is a count.

resolution V* and HRBench4K under the DeepEyes protocol (Zheng et al., 2025). Two cells (32B DYNAMO + RL on Chart and Table) are runs not yet completed.

Chart and Table reasoning. Using the gap-recovery fraction (DYNAMO − Direct)/(RL − Direct), DYNAMO recovers 65.7% and 99.4% of the Chart gap at 3B and 7B respectively, and exceeds RL at 32B (86.4 vs. 84.3). On Table, recovery is 67.5% / 91.4% / 46.8%. DYNAMO + RL further beats RL alone by 1–3 points wherever the row is reported, indicating that the two interventions compose additively.

High-resolution visual search. On V* and HRBench4K, DYNAMO alone matches or beats task-specific RL at every scale: it gains +6.3 / +1.0 / ±0.0 over RL on V* (3B / 7B / 32B) and +4.0 / +0.7 / +5.1 on HRBench4K. Adding DYNAMO on top of the RL-trained policy raises 7B V* to 92.7 and 7B HRBench4K to 81.3, the strongest configurations in those cells.

Cost. DYNAMO performs no weight updates: an evolution run uses 20% of each benchmark’s training split for $N=3$ passes and invokes the frozen VLM for solving, diagnosis, generation, and validation. The comparison therefore shows that across these three task families, DYNAMO recovers most of the RL gain on Chart/Table and matches or beats RL on high-resolution visual search, while composing additively with RL when both are available.

4.5 Ablation Studies

The *Skill Only* and *Tool Only* rows of Table 1 already isolate the contribution of each capability type, so we use the remaining ablation budget to study the *mechanisms* that drive those gains rather

Backbone	Variant	Chart	Table	V*	HRBench4K
Qwen2.5-VL-3B	Direct	45.3995	43.0921	71.20	63.62
	RL	66.95	56.25	78.01	66.50
	DYNAMO	59.56	51.97	84.29	70.50
	DYNAMO + RL	68.28	59.54	85.34	71.25
Qwen2.5-VL-7B	Direct	56.0532	56.9078	79.06	70.50
	RL	75.18	68.42	84.82	74.38
	DYNAMO	75.060	67.43	85.86	75.12
	DYNAMO + RL	76.88	69.41	92.67	81.25
Qwen2.5-VL-32B	Direct	81.3559	82.8947	81.68	75.62
	RL	84.3	84.3	88.48	72.88
	DYNAMO	86.44	83.55	88.48	78.00
	DYNAMO + RL			91.10	76.38

Table 5: **Self-evolution vs. task-specific RL across three task families.** Accuracy (%). DYNAMO alone matches or beats RL on all six high-resolution cells and recovers most of the Chart/Table gap; DYNAMO + RL is additive in nearly every cell. Blank cells are runs not yet completed.

than rerunning the same capability-type comparison. Specifically, we ask whether the improvement comes from (i) failure-specific capability generation as opposed to generic prompting, (ii) multi-modal failure diagnosis as opposed to text-only reflection, (iii) mastery-gated tool exposure as opposed to unconditional deployment, and (iv) persistent cross-case capability reuse as opposed to per-case retry. The full mechanism table, per-ablation discussion, and sensitivity curves are deferred to Appendix B for space.

5 Conclusion

DYNAMO is a training-free framework in which a frozen VLM agent builds and updates its own task-family-specific capability library from a small training subset. The agent inspects both correct and incorrect attempts on a sampled sub-training set, an evolution planner decides whether to evolve a reasoning skill, an executable visual tool, or both, and the best of multiple candidates is promoted by validation accuracy. A mastery skill paired with each tool gates its invocation at inference, and an online adaptation loop updates the library when the query distribution shifts.

Across four visual reasoning benchmarks and five VLM backbones, DYNAMO improves direct inference on all 20 model–benchmark cells. On the tool-mastery benchmark GTA, the same framework lifts step-level tool selection, argument, and instruction-following accuracy across every tested backbone. Compared with task-specific RL (VTool-R1, DeepEyes) on chart, table, and high-resolution perception tasks, DYNAMO matches or beats RL

on the high-resolution side and recovers most of the gap on the others, while composing additively when RL is also available. Together, these results suggest that a substantial part of task-specific VLM adaptation can be obtained without any gradient updates, by letting the agent shape its own capability library.

Limitations

Short evolution horizon. We report results at $N=3$ iterations; the evolution dynamics suggest gains plateau by $N=5$, but we have not evaluated whether longer runs (e.g. $N>10$) eventually overfit the training distribution.

Tool quality ceiling. Tool effectiveness is bounded by the backbone’s Python coding ability. For benchmarks requiring specialised computer-vision operations beyond standard libraries (e.g. learned feature extractors), the Generator may fail to produce a working tool, limiting the perceptual gains.

Benchmark scope. We evaluate four visual-reasoning benchmarks plus GTA and the VTool-R1 / DeepEyes splits. Broader coverage (video QA, 3D spatial reasoning, medical imaging) is needed before strong generality claims about the cognitive / perceptual decision rule.

API cost. Each training case incurs roughly 15K tokens for the combined attempt, diagnosis, generation, and validation pass. While negligible compared with RL training, this cost should be accounted for when comparing training-free methods.

References

- Anonymous. 2024. HRBench: A high-resolution visual benchmark for multimodal models (placeholder). [VERIFY] Replace with correct citation upon confirmation.
- Anonymous. 2025a. EvolveR: Iterative self-evolution of reasoning strategies (placeholder). [VERIFY] Replace with correct arXiv ID / venue upon confirmation.
- Anonymous. 2025b. PixelReasoner: Pixel-level reasoning for visual language models (placeholder). [VERIFY] Replace with final NeurIPS 2025 citation upon confirmation.
- Anonymous. 2025c. V-Thinker: Visual thinking for multimodal language models (placeholder). [VER-

587	IFY] Replace with correct arXiv ID upon confirma-	December 6-10, 2023, Findings of ACL, pages 6922–	644
588	tion.	6939. Association for Computational Linguistics.	645
589	Anonymous. 2025d. VTool-R1: Visual tool reasoning	Yujia Qin, Shihao Liang, Yining Ye, Kunlun Zhu, Lan	646
590	via reinforcement learning (placeholder). [VERIFY]	Yan, Yaxi Lu, Yankai Lin, Xin Cong, Xiangru Tang,	647
591	Replace with final ICLR 2026 citation upon confir-	Bill Qian, Sihan Zhao, Lauren Hong, Runchu Tian,	648
592	mation.	Ruobing Xie, Jie Zhou, Mark Gerstein, Dahai Li,	649
593	Tianle Cai, Xuezhi Wang, Tengyu Ma, Xinyun Chen,	Zhiyuan Liu, and Maosong Sun. 2024. ToolLLM:	650
594	and Denny Zhou. 2023. Large language models as	Facilitating large language models to master 16000+	651
595	tool makers . Preprint, arXiv:2305.17126.	real-world APIs . In <i>The Twelfth International Con-</i>	652
596	Minghao Chen, Yihang Li, Yanting Yang, Shiyu Yu,	<i>ference on Learning Representations, ICLR 2024,</i>	653
597	Binbin Lin, and Xiaofei He. 2024. AutoManual:	<i>Vienna, Austria, May 7-11, 2024</i> . OpenReview.net.	654
598	Constructing instruction manuals by LLM agents via	Haozhan Shen, Kangjia Zhao, Tiancheng Zhao,	655
599	interactive environmental learning . In <i>Advances in</i>	Ruochen Xu, Zilun Zhang, Mingwei Zhu, and	656
600	<i>Neural Information Processing Systems 38: Annual</i>	Jianwei Yin. 2024. ZoomEye: Enhancing multi-	657
601	<i>Conference on Neural Information Processing Sys-</i>	modal LLMs with human-like zooming capabili-	658
602	<i>tems 2024, NeurIPS 2024, Vancouver, BC, Canada,</i>	ties through tree-based image exploration . Preprint,	659
603	<i>December 10-15, 2024</i> .	arXiv:2411.16044. Accepted by EMNLP 2025.	660
604	Xingyu Fu, Minqian Liu, Zhengyuan Yang, John Cor-	Noah Shinn, Federico Cassano, Ashwin Gopinath,	661
605	ring, Yijuan Lu, Jianwei Yang, Dan Roth, Dinei Flo-	Karthik Narasimhan, and Shunyu Yao. 2023. Re-	662
606	rencio, and Cha Zhang. 2025. ReFocus: Visual edit-	flexion: language agents with verbal reinforcement	663
607	ing as a chain of thought for structured image under-	learning . In <i>Advances in Neural Information Pro-</i>	664
608	standing . Preprint, arXiv:2501.05452.	<i>cessing Systems 36: Annual Conference on Neural</i>	665
609	Tanmay Gupta and Aniruddha Kembhavi. 2023. Vi-	<i>Information Processing Systems 2023, NeurIPS 2023,</i>	666
610	sual programming: Compositional visual reasoning	<i>New Orleans, LA, USA, December 10 - 16, 2023</i> .	667
611	without training . In <i>IEEE/CVF Conference on Com-</i>	Dídac Surís, Sachit Menon, and Carl Vondrick. 2023.	668
612	<i>puter Vision and Pattern Recognition, CVPR 2023,</i>	ViperGPT: Visual inference via python execution for	669
613	<i>Vancouver, BC, Canada, June 17-24, 2023</i> , pages	reasoning . In <i>IEEE/CVF International Conference</i>	670
614	14953–14962. IEEE.	<i>on Computer Vision, ICCV 2023, Paris, France, Oc-</i>	671
615	Pan Lu, Hritik Bansal, Tony Xia, Jiacheng Liu, Chun-	<i>tober 1-6, 2023</i> , pages 11854–11864. IEEE.	672
616	yuan Li, Hannaneh Hajishirzi, Hao Cheng, Kai-	Guanzhi Wang, Yuqi Xie, Yunfan Jiang, Ajay Man-	673
617	Wei Chang, Michel Galley, and Jianfeng Gao. 2024.	dlekar, Chaowei Xiao, Yuke Zhu, Linxi Fan, and An-	674
618	MathVista: Evaluating mathematical reasoning of	ima Anandkumar. 2024a. Voyager: An open-ended	675
619	foundation models in visual contexts . In <i>The Twelfth</i>	embodied agent with large language models . <i>Trans.</i>	676
620	<i>International Conference on Learning Representa-</i>	<i>Mach. Learn. Res.</i> , 2024.	677
621	<i>tions, ICLR 2024, Vienna, Austria, May 7-11, 2024</i> .	Jize Wang, Zerun Ma, Yining Li, Songyang Zhang,	678
622	OpenReview.net.	Cailian Chen, Kai Chen, and Xinyi Le. 2024b. GTA:	679
623	Pan Lu, Baolin Peng, Hao Cheng, Michel Galley, Kai-	A benchmark for general tool agents . In <i>Advances in</i>	680
624	Wei Chang, Ying Nian Wu, Song-Chun Zhu, and	<i>Neural Information Processing Systems (NeurIPS),</i>	681
625	Jianfeng Gao. 2023. Chameleon: Plug-and-play com-	pages 75749–75790.	682
626	positional reasoning with large language models . In	Penghao Wu and Saining Xie. 2023. V*: Guided visual	683
627	<i>Advances in Neural Information Processing Systems</i>	search as a core mechanism in multimodal LLMs .	684
628	<i>36: Annual Conference on Neural Information Pro-</i>	Preprint, arXiv:2312.14135.	685
629	<i>cessing Systems 2023, NeurIPS 2023, New Orleans,</i>	Xintong Zhang, Zhi Gao, Bofei Zhang, Pengxiang Li,	686
630	<i>LA, USA, December 10 - 16, 2023</i> .	Xiaowen Zhang, Yang Liu, Tao Yuan, Yuwei Wu,	687
631	Ahmed Masry, Do Xuan Long, Jia Qing Tan, Shafiq R.	Yunde Jia, Song-Chun Zhu, and Qing Li. 2025. Adap-	688
632	Joty, and Enamul Hoque. 2022. ChartQA: A bench-	tive chain-of-focus reasoning via dynamic visual	689
633	mark for question answering about charts with visual	search and zooming for efficient VLMs . Preprint,	690
634	and logical reasoning . In <i>Findings of the Association</i>	arXiv:2505.15436.	691
635	<i>for Computational Linguistics: ACL 2022, Dublin,</i>	Andrew Zhao, Daniel Huang, Quentin Xu, Matthieu Lin,	692
636	<i>Ireland, May 22-27, 2022, Findings of ACL</i> , pages	Yong-Jin Liu, and Gao Huang. 2024. ExpeL: LLM	693
637	2263–2279. Association for Computational Linguis-	agents are experiential learners . In <i>Thirty-Eighth</i>	694
638	tics.	<i>AAAI Conference on Artificial Intelligence, AAAI</i>	695
639	Cheng Qian, Chi Han, Yi Ren Fung, Yujia Qin, Zhiyuan	<i>2024, February 20-27, 2024, Vancouver, Canada</i> .	696
640	Liu, and Heng Ji. 2023. CREATOR: Tool creation for	Ziwei Zheng, Michael Yang, Jack Hong, Chenxiao	697
641	disentangling abstract and concrete reasoning of large	Zhao, Guohai Xu, Le Yang, Chao Shen, and Xing	698
642	language models . In <i>Findings of the Association for</i>	Yu. 2025. DeepEyes: Incentivizing “thinking	
	<i>Computational Linguistics: EMNLP 2023, Singapore,</i>		

with images” via reinforcement learning. *Preprint*,
arXiv:2505.14362. Accepted by ICLR 2026.

A Implementation Details

Algorithm. Algorithm 1 gives the full DYNAMO evolution procedure described in Section 3.2.

Algorithm 1: DYNAMO evolution loop.	
Input: $\mathcal{D}_{\text{train}}$, frozen VLM π_θ , iterations N , candidates M	
Output: Capability set $\mathcal{C} = (\mathcal{S}, \mathcal{T})$	
1	$\mathcal{C} \leftarrow (\emptyset, \emptyset);$
2	for $n = 1$ to N do
3	Solve $\mathcal{D}_{\text{train}}$ with $f_{\mathcal{C}}$; collect attempts;
4	$\mathcal{D}_{\text{sub}} \leftarrow$ sample from $\mathcal{D}_{\text{train}}$;
5	$a, r \leftarrow$
	ANALYZERDECIDER($\mathcal{D}_{\text{sub}}; \pi_\theta$);
6	$\{(s_m, t_m)\}_{m=1}^M \leftarrow$
	GENERATOR($a, r; \pi_\theta$); // s_m
	paired with t_m acts as its
	mastery SOP
7	$m^* \leftarrow$
	$\arg \max_m \text{Acc}(f_{\mathcal{C} \cup \{(s_m, t_m)\}}; \mathcal{D}_{\text{train}});$
8	if accuracy improves then
9	$\mathcal{C} \leftarrow \mathcal{C} \cup \{(s_{m^*}, t_{m^*})\};$
10	end
11	end
12	return $\mathcal{C}$;

Role prompts. DYNAMO uses three role-specific system prompts for the same base VLM: *Solver* (standard QA prompt with capability context injected), *AnalyzerDecider* (root-cause analysis and decision), and *Generator* (skill or tool generation). All prompts are provided in full at [\[anonymised repository\]](#).

Skill retrieval. At inference time, the Solver retrieves the top- $K=3$ skills from $\mathcal{S}$ using BM25 cosine similarity over skill titles and When-to-Use fields. Retrieved skill text is appended to the system prompt.

Tool invocation. The Solver inspects each tool’s mastery SOP to decide whether to invoke it on the current input. If the input matches a tool’s supported patterns, the tool is applied and its output image(s) are appended to the multimodal context. Tool execution is sandboxed with a 30-second timeout and resource limits.

Hyperparameters. Table 6 lists all hyperparameter settings.

Table 6: Hyperparameter settings.

Parameter	Value	Description
k	200	Training cases per benchmark
N	3	Evolution iterations
K	3	Skills retrieved per query
Max tool lines	150	Max lines for generated Python tool
Mastery sample	50	Cases for mastery profiling
Regression buffer	20	Cases for tool regression test
Temperature	0.7	Generator temperature
Max tokens (gen.)	2048	Max tokens for capability generation

Compute. Experiments use the frozen VLM backbones listed in Section 4.1. Total API calls per benchmark: $\leq k \times N \times 3 = 1,800$ calls (solver + analyzer + generator, upper bound). Estimated token usage: $\sim 3\text{M}$ tokens per benchmark. No GPU training is performed.

B Mechanism Ablations

Experiment I in the main paper already isolates the contribution of each capability *type*: the *None*, *Skill Only*, *Tool Only*, and *Skill + Tool* rows of Table 1 answer whether skills, tools, or their combination is responsible for the observed gains. The ablations in this appendix test a complementary question: when the full capability library is used, which *mechanisms* inside the evolution loop are doing the work? Each variant below removes one design decision—failure-specific capability generation, multimodal failure diagnosis, mastery-gated tool exposure, or cross-case persistence—while holding the rest of the system fixed. The unifying claim under test is that the gains come from *failure-driven capability evolution*: the agent visually diagnoses failures, generates persistent skills and tools, validates them, and learns when to deploy them through mastery. Removing any of these mechanisms either reduces improvement or makes tool use less reliable.

Setup. All variants use the same protocol as Experiment I (frozen backbone; $k=200$ training cases; $N=3$ evolution iterations) and rebuild the capability library from scratch so that any deficit reflects the ablated mechanism rather than a carried-over artifact. We report results on a representative backbone (we use GPT-5.4 in the main mechanism table, with GPT-4o and Doubao-Seed-2.0 entries reported below as robustness checks where available). The two main benchmarks are chosen to stress differ-

Variant	Ablated mechanism	ChartQA	HRBench	MathVista	V*
Full DYNAMO	none (full evolution loop)	0.859	0.850	0.786	0.874
Direct (reference)	no capability library; zero-shot VLM	0.797	0.747	0.759	0.775
Generic Skill	no failure-specific evolution; hand-neutral task prompt only	0.815	0.772	0.769	0.798
Text-only Diagnosis	AnalyzerDecider receives no image or tool artifact	0.844	0.802	0.778	0.815
No Mastery	tools exposed without applicability SOP	0.838	0.823	0.754	0.836
No Persistence	reflection/capability not stored across cases	0.819	0.819	0.770	0.828

Table 7: **Mechanism ablations.** Each row removes one mechanism from DYNAMO while holding all others fixed (GPT-5.4 backbone; $k=200$, $N=3$). The ablated dimensions are *mechanisms* (how capabilities are generated, diagnosed, gated, and reused) rather than *capability types*, which are already ablated by the *Skill Only / Tool Only* rows of Table 1. The Full row reproduces the Skill + Tool entry of Table 1, and the Direct row reproduces the None (Base Agent) entry, both for reference. The four ablation rows are *layout placeholders* chosen to illustrate the expected qualitative pattern and must be replaced with measured values before submission.

ent bottlenecks: **ChartQA** as a cognitive / procedural test where persistent reasoning skills are expected to help, and **HRBench** as a perceptual / high-resolution test where visual failure analysis and mastery-gated tools matter most. **MathVista** and **V*** entries are reported as a secondary check.

Main mechanism table. Table 7 reports the five variants on the two primary benchmarks. Each row maps to one specific design decision in Section 3, and the corresponding subsections below describe the variant, its motivation, and the expected reading of the numbers. The Full DYNAMO row reproduces the Skill + Tool entry from Table 1 for reference; the other rows are not subsumed by any row of that table.

Reading the table. Across the four mechanism variants we expect a consistent qualitative pattern: Generic Skill should recover only a small fraction of the Full gain (it captures shallow prompt advice but not failure-specific procedures and triggers); Text-only Diagnosis should hurt most on the perceptual benchmarks where the AnalyzerDecider must inspect crops and processed artifacts; No Mastery should leave or even increase tool usage but reduce reliability by invoking tools on cases that direct reasoning already solves; and No Persistence should preserve per-case retry benefits but lose held-out gains that depend on retrieving capabilities discovered on different training cases.

B.1 Generic Skill vs. Evolved Capability

Question. Are the gains caused by failure-specific evolution, or would a generic task-aware

prompt do the same thing?

Variants. We compare three configurations on the same backbone and the same evaluation protocol. *Direct* is the zero-shot VLM with no persistent capability and no special prompt. *Generic Skill* replaces the evolved capability library with a single hand-neutral task prompt that gives general task advice not derived from any failure analysis: for ChartQA, the prompt instructs the model to “read the chart carefully, identify the relevant values, and reason step by step”; for HRBench, “inspect local details before answering and verify the target object before committing to an answer.” *Evolved Skill/Tool* is the frozen capability library produced by the standard evolution loop.

Expected interpretation. A pure prompt-engineering account of the main results would predict that Generic Skill closes most of the gap to the evolved library, because the evolved skills are themselves expressed as natural-language SOPs. The hypothesis under test is the opposite: that Generic Skill captures only shallow task advice, while the evolved capabilities encode failure-specific procedures, applicability triggers, and pitfalls that transfer beyond the triggering examples. Concretely, if Generic Skill improves over Direct by a small margin but remains well below Full, the gap is attributable to mechanism-level evolution rather than to the presence of any task-aware prompt. On ChartQA in particular, evolved skills frequently encode chart-type-conditioned decomposition patterns (e.g., “for stacked bars, read the upper and lower boundaries of the target

segment before subtracting”) that a generic chart prompt does not specify.

Writing hook. Generic task instructions provide a useful but incomplete baseline: they tell the model to be careful, but they do not identify which recurring failure modes were observed during training or when a visual tool should be invoked.

B.2 Text-only Diagnosis

Question. Does the AnalyzerDecider need to see images and intermediate tool artifacts, or is a text-only reflection sufficient?

Variants. *Full Diagnosis* is the standard configuration in which the AnalyzerDecider receives the question, the failed answer, the ground-truth answer, the full reasoning trace, the original image, and any intermediate visual artifacts produced by tools invoked during the failed attempt. *Text-only Diagnosis* restricts the AnalyzerDecider’s inputs to question, failed answer, ground truth, and trace; no original image and no generated, cropped, or processed visual artifacts are passed. The rest of the loop is unchanged, so any difference reflects only the diagnostic channel.

Expected interpretation. Text-only reflection can describe *that* an answer was wrong and can speculate about plausible cognitive causes, but it cannot determine whether a crop was misaligned, whether a zoomed region missed the target, whether a generated visual artifact degraded the evidence, or whether the original image already contained the information needed. We therefore expect the gap between Full and Text-only Diagnosis to be largest on HRBench and V*, where correctly diagnosing a perceptual failure requires inspecting the image. On ChartQA, the gap should be smaller because many failures there are procedural rather than perceptual; a small or negligible gap on ChartQA is consistent with, rather than evidence against, the visual-diagnosis claim. The asymmetric pattern itself supports the cognitive/perceptual distinction that motivates the dual-capability design.

Writing hook. Text-only reflection can describe that an answer was wrong, but it cannot tell whether a crop was misaligned, a zoomed region missed the target, or a generated visual artifact degraded the evidence shown to the next solver call. 920

B.3 No Mastery / Unconditional Tool Exposure

Question. Does progressive mastery matter, or can the agent simply expose every validated tool whenever one is available?

Variants. *Full* validates each generated tool and gates its exposure on a learned mastery SOP that describes the inputs and contexts in which the tool tends to improve over direct reasoning. *No Mastery* keeps validation but removes the mastery SOP: any tool that passes basic syntax and origin-case checks is exposed to the Solver on every case, with at most a coarse dataset-level rule. As an optional bound, *All Tools Exposed* exposes every tool on every case regardless of validation; we treat this as a stress test rather than a main row because it is rarely used in practice.

Expected interpretation. No Mastery is expected to increase tool-usage rate but reduce accuracy or increase variance. This is the central motivation for the mastery design: many generated tools are useful only in a narrow input region, and indiscriminate exposure causes the Solver to invoke them on cases where direct reasoning was already sufficient, introducing unnecessary visual artifacts or intermediate steps. The MathVista cell is informative here: prior runs in the main paper show that naive tool exposure can hurt MathVista more than it helps, and we expect the No Mastery variant to reproduce that pattern. The Full vs. No Mastery comparison therefore tests whether mastery improves *precision* rather than just increasing tool-use rate.

Writing hook. Mastery improves precision rather than simply increasing tool-use rate. The unconditional variant often calls tools on cases where direct reasoning is already sufficient, introducing visual artifacts or unnecessary intermediate steps.

B.4 No Persistence / Per-case Reflection

Question. Does storing reusable capabilities across cases matter, or is it enough to reflect and retry on each failed case independently?

Variants. *Persistent Library* is the standard configuration: promoted skills and tools are stored in *C* and retrieved on later training and evaluation cases. *Per-case Reflection* allows the agent to reflect on a failed attempt and immediately retry the *same* case using the generated capability, but discards the ca-

pability afterwards so it is unavailable to any future case. In both variants, the AnalyzerDecider and Generator are otherwise identical; only the storage and retrieval interface is removed.

Expected interpretation. Per-case Reflection should help on the triggering case but provide little benefit on held-out cases, because each failure is repaired in isolation. The gap between Persistent Library and Per-case Reflection on held-out accuracy therefore measures the value of *cross-case capability reuse*, not the value of reflection per se. ChartQA is the cleanest test: failures on chart questions form recurring families (stacked-bar boundary reading, multi-series selection, percentage versus absolute value), and the same generated skill should fix many held-out cases of the same family. If Persistent Library substantially outperforms Per-case Reflection on held-out ChartQA accuracy, the system’s advantage comes from accumulating reusable capabilities rather than from per-case retry alone.

Writing hook. The goal of evolution is not to repair a single failed attempt, but to convert that failure into a reusable capability that future cases can retrieve.

B.5 Optional: Validation and Regression Filter

This ablation tests whether the multi-stage validation pipeline that gates capability promotion is necessary, rather than which mechanism produces the capability in the first place. *Full Validation* requires a generated tool to pass syntax checking, origin-case improvement (the tool must fix the triggering case), and a regression check over a buffer of previously-correct cases before mastery profiling begins. *Origin-only* accepts a tool as soon as it fixes the triggering case, skipping the regression buffer. *No Regression* keeps the syntax and origin checks but disables the prior-correct buffer.

Origin-only is expected to produce higher training-set recovery but worse held-out accuracy because it admits tools that overfit the triggering case or harm visually similar but semantically different inputs. HRBench and V* are the most informative benchmarks here because generated visual tools (cropping, zooming, masking) can easily overfit to one image region. We report this ablation as optional because the regression-buffer cost is the same order as one extra validation pass, and its

main value is in preventing failure modes that appear only with larger or longer evolution runs than the main experiments use.

B.6 Sensitivity to Training Subset Size and Iteration Count

The main paper fixes $k=200$ training cases and $N=3$ evolution iterations. The sensitivity analysis in Appendix C (Figure 3) sweeps $k \in \{10, 25, 50, 100, 200\}$ at $N=3$ and reports ChartQA accuracy. We expect a monotone but quickly saturating curve: a small k already exposes the most frequent failure families, and adding more training cases mostly contributes redundant capabilities that the persistence and validation filters merge or reject. A similar saturation is expected for N : $N=1$ misses failures that only appear after earlier capabilities reshape the trajectory, while $N=3$ and $N=5$ tend to produce comparable libraries because most remaining failures fall outside the in-distribution training subset. We treat the sensitivity sweep as a robustness check rather than a tuning study; the operating point is chosen to fit the per-benchmark budget rather than to maximize accuracy.

B.7 Summary

Read together, the four mechanism ablations are designed to triangulate the same claim from four different angles. Generic Skill rules out a pure prompt-engineering explanation; Text-only Diagnosis rules out a pure text-reflection explanation and supports the visual-diagnosis channel; No Mastery rules out an “expose all generated tools” explanation and supports applicability-bounded tool use; and No Persistence rules out a per-case retry explanation and supports cross-case capability accumulation. The expected pattern is that removing any one mechanism reduces held-out accuracy noticeably, but no single ablation collapses the system, which is consistent with the design view that the four mechanisms implement complementary aspects of capability evolution rather than redundant copies of the same underlying lever.

Practical diagnostic. The mechanism ablations also suggest a practical rule of thumb for new task families. If the bottleneck is procedural—visual evidence is present but the reasoning is weak—skill evolution is the lower-cost first choice; if a different view of the image is required before the evidence becomes accessible, tool evolution is necessary. A

practitioner targeting a new benchmark can use a small pilot run (e.g. 5% of the training split with $N=1$) to inspect which kinds of capabilities are generated and whether they transfer beyond the triggering examples.

C Additional Results

C.1 Training Set Size Sensitivity



Figure 3: Placeholder for ChartQA accuracy vs. training set size $k \in \{10, 25, 50, 100, 200\}$ at $N=3$ iterations.

Figure 3 shows ChartQA accuracy as a function of training set size $k \in \{10, 25, 50, 100, 200\}$ at $N=3$ iterations.

C.2 Per-Category Analysis on ChartQA

D Qualitative Analysis of Evolved Capabilities

This appendix gives three qualitative case studies of evolved capabilities. The examples are intended to make the generated artifacts concrete, not to introduce new quantitative results. Each case should be instantiated with the corresponding logged failure, generated skill/tool file, visual artifact, and validation record before submission. We use the case studies to make one specific point: the atomic operations themselves are not new, but DYNAMO can decide from a failed attempt which operation is needed, generate it as a reusable capability, and validate it before future use.

D.1 Case A: ReFocus-Style Skill and Tool for Structured Images

ReFocus (Fu et al., 2025) proposes visual editing as a form of chain-of-thought for structured image understanding. It asks an MLLM to generate Python code that edits the image, including drawing boxes, highlighting relevant regions, and masking distractors. A DYNAMO case belongs in this category when the generated capability explicitly changes the visual input shown to the solver.

Triggering failure. The triggering case should be a structured image example, such as a chart or

table, where the base agent attends to the wrong visual element or confuses visually similar elements. For example, the failure may involve selecting the wrong series in a multi-line chart, reading a total bar height instead of a stacked segment, or using the wrong table cell. [Insert benchmark, case id, question, ground-truth answer, and base answer.]

Generated skill. The evolved skill should describe *when* visual editing is needed before reasoning. A typical ReFocus-style skill has the following structure:

```
## Skill: Edit-Then-Read for Structured Images

When-to-Use:
Use when the question refers to a specific chart/table element
that is visually close to distractors, such as one series
among
many lines, one segment in a stacked bar, or one table
row/column.

Strategy:
1. Identify the visual entity named by the question.
2. Before computing the answer, create an edited view that
   marks
   only the relevant region and suppresses likely distractors.
3. Re-read the edited image and extract the required value.
4. Perform the requested arithmetic or comparison.

Common Pitfalls:
- Do not answer from the unedited image if multiple marks are
  visually similar.
- Do not treat the visual edit as adding new information; it
  only
  exposes the evidence already present in the image.
```

Generated tool. The paired tool is a small image-editing program. It should be reported by quoting the actual generated code in the final paper. A representative form is:

```
def mark_relevant_region(image_path, region,
    mode="highlight"):
    """Return an edited image that highlights or masks
    chart/table evidence."""
    from PIL import Image, ImageDraw
    img = Image.open(image_path).convert("RGB")
    draw = ImageDraw.Draw(img, "RGBA")
    x1, y1, x2, y2 = region
    if mode == "highlight":
        draw.rectangle([x1, y1, x2, y2], outline=(255, 0, 0,
            255), width=6)
        draw.rectangle([x1, y1, x2, y2], fill=(255, 255, 0,
            60))
    elif mode == "mask_distractors":
        # final paper should quote the actual generated
        # masking logic
        pass
    out_path = image_path.replace(".png", "_marked.png")
    img.save(out_path)
    return out_path
```

Why this matches ReFocus. Like ReFocus, the capability uses visual editing as an intermediate reasoning step: the model does not merely write a textual explanation, but creates an edited visual artifact that changes the evidence shown to the next solver call. The difference is that ReFocus defines visual editing as the reasoning interface, whereas DYNAMO produces this edit skill/tool only after a failure analysis decides that the bottleneck is visual disambiguation.

Case study	Generated capability	Prior motif	Claim supported
Case A: Structured-image editing	A skill that instructs the solver to isolate the relevant chart/table structure, plus a tool that highlights, boxes, or masks visual evidence.	ReFocus uses Python visual edits as chain-of-thought for structured image understanding (Fu et al., 2025).	DYNAMO can generate a reusable visual-editing capability from failure, rather than relying on a fixed visual-editing interface.
Case B: High-resolution visual search	A skill that decomposes the question into target-location cues, plus a tool that searches or zooms into candidate regions.	ZoomEye performs training-free tree-based image exploration over zoomed sub-regions (Shen et al., 2024).	DYNAMO can generate a local coarse-to-fine search capability without being given ZoomEye as a predefined module.
Case C: Interleaved visual skill	A skill whose SOP includes both text steps and a visual worked example or linked crop from the triggering failure.	Visual-thought methods show that intermediate visual artifacts can guide structured reasoning.	DYNAMO can store procedural and visual evidence together as a retrieved capability.

Table 8: **Three qualitative forms of evolved capability.** The case studies compare generated artifacts to prior visual-reasoning motifs. They do not claim that DYNAMO reproduces the full prior algorithms.

Evidence to insert. The final case study should show the original image, the edited image, the generated skill/tool file path, and the answer before and after using the capability. **[Insert validation result and retry answer.]**

D.2 Case B: ZoomEye-Style Skill and Tool for High-Resolution Search

ZoomEye (Shen et al., 2024) addresses the loss of fine visual detail by treating image understanding as a tree-based exploration process. The full image is the root node, zoomed image regions are child nodes, and the model searches over promising regions until it finds the visual evidence needed to answer. A DYNAMO case belongs in this category when the generated capability performs region selection, cropping, zooming, or coarse-to-fine search.

Triggering failure. The triggering case should involve a high-resolution image where the base agent knows the object or region type but cannot inspect it accurately from the full image. Examples include small text, distant signs, number plates, small object attributes, or a target object occupying a small fraction of the image. **[Insert benchmark, case id, question, ground-truth answer, and base answer.]**

Generated skill. The evolved skill should make the search procedure explicit:

```
## Skill: Coarse-to-Fine Region Search
```

```
When-to-Use:
Use when the question asks about a small object, distant text,
fine-grained attribute, or localized region in a
high-resolution image.
```

```
Strategy:
```

1. Parse the question for target object, attribute, and spatial cues.
2. Inspect the full image only to estimate candidate regions.
3. Divide the image into coarse tiles and score each tile for relevance.
4. Zoom into the best candidate region and answer from the enlarged view.
5. If the crop is ambiguous, repeat with a finer crop around the target.

Common Pitfalls:

- Do not answer from the globally downsampled image when the requested evidence is small.
- Do not crop solely by image center; use question-specific cues.

Generated tool. The paired tool implements a local search/zoom operation. The final version should quote the actual generated code; a representative form is:

```
def coarse_to_fine_zoom(image_path, target_description,
    grid_size=4):
    """Return candidate high-resolution crops for a
    target described in text."""
    from PIL import Image
    img = Image.open(image_path).convert("RGB")
    w, h = img.size
    crops = []
    for gy in range(grid_size):
        for gx in range(grid_size):
            box = (gx*w//grid_size, gy*h//grid_size,
                (gx+1)*w//grid_size, (gy+1)*h//grid_size)
            crop = img.crop(box)
            out = image_path.replace(".png",
                f"_tile_{gx}_{gy}.png")
            crop.save(out)
            crops.append((out, box))
    return crops
```

Why this matches ZoomEye. The similarity is the coarse-to-fine visual search motif. Both methods avoid forcing the VLM to answer from a single fixed-resolution full image. The difference is scope: ZoomEye is a general tree-search algorithm, while the DYNAMO capability is generated from a particular failure and may implement only the subset of zoom/search behavior needed for that task family.

Evidence to insert. The final case study should

include the full image, the candidate tile or crop selected by the tool, and the before/after answer. **[Insert selected crop, validation result, and retry answer.]**

D.3 Case C: Interleaved Visual Skill

The third case is different from the previous two because the main artifact is not only an executable program. An interleaved visual skill is a retrieved Markdown skill that contains both textual reasoning instructions and visual evidence, such as an embedded crop, annotated chart, or linked processed image from the triggering failure.

Triggering failure. This case should involve a problem where a text-only SOP is underspecified because the visual pattern itself matters. A common example is chart reasoning: the agent may know that it should subtract segment boundaries in a stacked bar chart, but still confuse which boundary belongs to the target segment. **[Insert benchmark, case id, question, ground-truth answer, and base answer.]**

Generated skill. The generated skill should be quoted directly if available. A representative interleaved form is:

```
## Skill: Stacked-Bar Segment Reading with Visual Anchor

When-to-Use:
Use when the question asks for a value, difference, or
comparison
involving a segment of a stacked bar chart.

Strategy:
1. Locate the target bar using the x-axis label.
2. Locate the target segment using the legend colour.
3. Read the lower and upper boundaries of the segment.
4. Subtract lower from upper to obtain the segment value.
5. Verify the result against the visual anchor below.

Visual Anchor:
[embedded crop or linked image showing the annotated segment
boundaries]

Common Pitfalls:
- Do not read the total bar height when the question asks for
one segment.
- Do not use a nearby segment with a visually similar colour.
```

Optional generated tool. In this case, the paired tool may be simple: it produces the crop or annotation that is embedded in the skill.

```
def make_visual_anchor(image_path, crop_box, annotations):
    """Create an annotated crop used as the visual anchor in
    a skill."""
    from PIL import Image, ImageDraw
    img = Image.open(image_path).convert("RGB")
    crop = img.crop(crop_box)
    draw = ImageDraw.Draw(crop)
    for ann in annotations:
        draw.rectangle(ann["box"], outline=ann.get("color",
"red"), width=4)
        if "label" in ann:
            draw.text((ann["box"][0], ann["box"][1]),
ann["label"],
                    fill=ann.get("color", "red"))
    out_path = image_path.replace(".png",
    "_visual_anchor.png")
```

```
crop.save(out_path)
return out_path
```

Why this is interleaved. The retrieved capability is not purely textual and not purely executable. The skill tells the agent how to reason, while the visual anchor shows what the relevant visual configuration looks like. This is related to visual-thought methods, but the visual artifact is stored persistently inside the capability library and retrieved on later cases.

Evidence to insert. The final case study should include the generated Markdown skill, the visual anchor image, and a before/after answer showing how the retrieved interleaved skill changes the solver's behavior. **[Insert validation result and retry answer.]**




Figure 4: **Placeholder for a concrete before/after case study.** The final version should insert the original image, the failed base answer, the generated skill/tool artifact, the processed visual output if any, and the validated retry answer.

E Case Study Artifact Checklist

To avoid conflating qualitative analysis with unreported experiments, we instantiate each case study only from artifacts produced by the evolution loop. For each of the three cases above, the final appendix should include:

- the triggering failure record: benchmark, example id, question, ground-truth answer, and base-agent answer;
- the generated skill file, quoted or summarized without changing its operative steps;
- the generated tool file when one exists, including the exact operation it performs;

- the visual artifact produced by the tool, such as a marked chart, selected crop, or embedded visual anchor;
- the validation record showing whether the capability was accepted and whether the retry answer changed.

If a case lacks one of these artifacts, we mark the missing field explicitly rather than inferring it from the intended behavior. This keeps the case studies as evidence for the form of the evolved capabilities, while the quantitative claims remain grounded in the tables in Section 4.

F Reproducibility Statement

The code, evolved capability files (learned/ directory), and all evaluation scripts will be released in an anonymised repository upon acceptance. The experiments use publicly available benchmark datasets and the VLM backbones described in Section 4.1. Evolved capability sets (*.skill.md and *.tool.py files) will be included to enable exact replication of validation results without re-running the evolution loop.

1322